

On the Computation of Necessary and Sufficient Explanations

Adnan Darwiche, Chunxi Ji

Computer Science Department
University of California, Los Angeles
darwiche@cs.ucla.edu, jich@cs.ucla.edu

Abstract

For Boolean classifiers, the *complete reason* behind a decision is a Boolean formula that characterizes why the decision was made. This recently introduced notion has a number of applications, which include generating explanations, detecting decision bias and evaluating counterfactual queries. Prime *implicants* of the complete reason are known as *sufficient reasons* for the decision and they correspond to what is known as PI explanations and abductive explanations. In this paper, we refer to the prime *implicants* of a complete reason as *necessary reasons* for the decision. We justify this terminology semantically and show that necessary reasons correspond to what is known as contrastive explanations. We also study the computation of complete reasons for multi-class decision trees and graphs with nominal and numeric features for which we derive efficient, closed-form complete reasons. We further investigate the computation of shortest necessary and sufficient reasons for a broad class of complete reasons, which include the derived closed forms and the complete reasons for Sentential Decision Diagrams (SDDs). We provide an algorithm which can enumerate their shortest necessary reasons in output polynomial time. Enumerating shortest sufficient reasons for this class of complete reasons is hard even for a single reason. For this problem, we provide an algorithm that appears to be quite efficient as we show empirically.

Introduction

Reasoning about the behavior of AI systems has been receiving significant attention recently, particularly the decisions made by machine learning classifiers. Some methods operate directly on classifiers, e.g., (Ribeiro, Singh, and Guestrin 2016, 2018) while others operate on symbolic encodings of their input-output behavior, e.g., (Narodytska et al. 2018; Ignatiev, Narodytska, and Marques-Silva 2019a) which may be compiled into tractable circuits (Chan and Darwiche 2003; Shih, Choi, and Darwiche 2018b, 2019; Shi et al. 2020; Audemard, Koriche, and Marquis 2020; Huang et al. 2021a). When explaining decisions, the notion of a *sufficient reason* has been well investigated. This is a minimal subset of an instance that is sufficient to trigger the decision and can therefore be used to explain why it was made. Sufficient reasons were introduced in (Shih, Choi, and Darwiche 2018b) under the name of *PI explanations* and

later referred to as *abductive explanations* (Ignatiev, Narodytska, and Marques-Silva 2019a).¹ A related notion is *contrastive explanations* (Lipton 1990; Miller 2019) as formalized in (Ignatiev et al. 2020) which we discuss at length later.

(Darwiche and Hirth 2020) introduced the *complete reason* for a decision as a Boolean formula that characterizes why a decision was made, and showed how it can be used to gather insights about the decision. This includes generating explanations, determining decision bias and evaluating counterfactual queries. For example, it was shown that sufficient reasons correspond to the *prime implicants* of the complete reason. Hence, if one has access to the complete reason behind a decision, then one can abstract the computation of sufficient reasons away from the classifier and its encoding or compilation. Consider a classifier for admitting applicants to an academic program based on five Boolean features (Darwiche and Hirth 2020): passing the entrance exam (E), being a first time applicant (F), having good grades (G), having work experience (W) and coming from a rich hometown (R). The positive instances of this classifier are specified by the following Boolean formula: $\Delta = (e \vee g) \wedge (e \vee r) \wedge (e \vee w) \wedge (f \vee r) \wedge (\bar{f} \vee g \vee w)$. Luna (δ) passed the entrance exam, has good grades and work experience, comes from a rich hometown but is not a first time applicant ($\delta = e, f, g, r, w$). The classifier will admit Luna. The complete reason for this decision is: $\Gamma = (e \vee g) \wedge (e \vee w) \wedge (r) \wedge (\bar{f} \vee g \vee w)$. There are four prime implicants of Γ : (e, g, r) , (e, r, w) , (e, \bar{f}, r) and (g, r, w) . Each is a minimal subset of instance δ which is sufficient to trigger the admit decision. Even though the number of sufficient reasons may be exponential, the complete reason can be compact and computed in linear time if the classifier is represented using a suitable form (Darwiche and Hirth 2020). Further insights can be obtained about a decision by analyzing its complete reason. For example, the decision on Luna is *biased* as it would be different if she did not come from a rich hometown. In that case, she would be denied admission *because* she does not come from a rich hometown and is not a first time applicant as this would be the only sufficient reason for rejection. These conclusions

¹See, e.g., (Choi, Xue, and Darwiche 2012; Ribeiro, Singh, and Guestrin 2018; Wang, Khosravi, and den Broeck 2021) for some approaches that can be viewed as approximating sufficient reasons and (Ignatiev, Narodytska, and Marques-Silva 2019b) for a study of the quality of some of these approximations.

can be derived by operating directly, and efficiently, on the complete reason as shown in (Darwiche and Hirth 2020).

More recently, (Darwiche and Marquis 2021) introduced the notion of *universal literal quantification* to Boolean logic and used it to formulate complete reasons. According to this formulation, we can obtain the above complete reason Γ by computing $\forall e, f, g, r, w \cdot \Delta$, to be explained later. We will base our treatment on this formulation while operating in a discrete instead of a Boolean setting. The conclusion section in (Darwiche and Marquis 2021) proposed a generalization of universal literal quantification to discrete variables but without further discussion. We will adopt this definition, study it further and exploit it to derive efficient, closed-form complete reasons for multi-class decision trees and graphs with nominal (discrete) and numeric (continuous) features. We will show that the obtained complete reasons belong to a particular logical form that arise when explaining the decisions of a broader class of classifiers. We will further show that the *prime implicants* of complete reasons correspond to contrastive explanations, which will provide further insights into the semantics and utility of these explanations. We will refer to these prime implicants as *necessary reasons* for the decision and semantically justify this terminology. We will then propose an output polynomial algorithm for computing the shortest necessary reasons of the identified class of complete reasons. We will finally show that computing shortest sufficient reasons is hard for this class of complete reasons and propose an algorithm for computing them which appears to be quite efficient based on an empirical evaluation.

Proofs of all results and the code used for the empirical evaluation can be found in the supplementary material.

Syntax and Semantics of Discrete Formulas

We start by defining the syntax and semantics of discrete formulas which we use to capture classifiers with discrete features. The treatment in this section is largely classical and provides obvious generalizations of what is known on Boolean logic. But we spell it out so we can provide a formal treatment of our upcoming results, especially that we sometimes depart from what may be customary.

For a discrete variable X with values x_1, \dots, x_n , we will call $X = x_i$ a *state* for variable X . A *discrete formula* is defined over a set of discrete variables as follows. Every state or constant (\top, \perp) is a discrete formula. If α and β are discrete formulas, then $\neg\alpha$, $\alpha \vee \beta$ and $\alpha \wedge \beta$ are discrete formulas. A *positive literal* is a state $X = x_i$ typically denoted by x_i . A *negative literal* is a negated state $\neg(X = x_i)$, typically denoted by \bar{x}_i . A negative literal will also be called a state if the variable has only two values. A *clause* is a disjunction of literals with at most one literal per variable. A *term* is a conjunction of literals with at most one literal per variable. A *CNF* is a conjunction of clauses. A *DNF* is a disjunction of terms. An *NNF* is defined as follows. Constants and literals are NNFs. If α and β are NNFs, then $\alpha \vee \beta$ and $\alpha \wedge \beta$ are NNFs (hence, conjunctions and disjunctions cannot be negated). An NNF is \vee -*decomposable* iff for each disjunction $\bigvee_i \alpha_i$ in the NNF, the disjuncts α_i do not share variables. An NNF is \wedge -*decomposable* iff for each conjunction $\bigwedge_i \alpha_i$ in the NNF, the conjuncts α_i do not share variables.

An NNF is *positive* iff it contains only positive literals. Any NNF can be made positive by replacing negative literals \bar{x}_i with $\bigvee_{j \neq i} x_j$. An NNF is *monotone* iff it is positive and does not contain distinct states x_i and x_j for any variable X .

A *positive term* contains only positive literals (i.e., states). The *conditioning* of discrete formula Δ on positive term γ is denoted $\Delta|\gamma$ and obtained as follows. For each state $x_i \in \gamma$, replace the occurrences of x_i with \top and the occurrences of x_j , $j \neq i$, with \perp . The formula $\Delta|x_i$ does not mention variable X . An *instance* is a positive term which contains precisely one state for each variable. If we condition a discrete formula on an instance, we get a Boolean formula that does not mention any variables (evaluates to true or false).

The semantics of discrete formulas is symmetric to the semantics of Boolean formulas, except that the notion of a *world* (truth assignment) is now defined as a function that maps each discrete variable to one of its states (a world corresponds to an instance). A world ω *satisfies* a discrete formula α , written $\omega \models \alpha$, precisely when $\alpha|\omega$ evaluates to true. In this case, we say that world ω is a *model* of formula α . Notions such as satisfiability, validity, implication and equivalence can now be defined for discrete formulas as in Boolean logic. For example, formula α implies formula β , written $\alpha \models \beta$, iff every model of α is a model of β . We next define the notions of implicants and implicants. An *implicant* of a discrete formula Δ is a term δ such that $\delta \models \Delta$. The implicant is *prime* iff no other implicant δ^* is such that $\delta^* \subset \delta$. An *implicate* is a clause δ such that $\Delta \models \delta$. The implicate is *prime* iff no other implicate δ^* is such that $\delta^* \subset \delta$.

Our treatment will represent a classifier with discrete features and multiple classes c_1, \dots, c_n by a set of mutually exclusive and exhaustive discrete formulas $\Delta^1, \dots, \Delta^n$, where the models of formula Δ^i capture the instances in class c_i . That is, instance δ is in class c_i iff $\delta \models \Delta^i$. We refer to each Δ^i as a *class formula*. When $\delta \models \Delta^i$, we say that instance δ is *decided positively* by Δ^i . The complete reason for this decision will then be the formula $\forall \delta \cdot \Delta^i$. The next section will explain what $\forall \delta$ is and how to compute it efficiently. In the upcoming discussion, we may use the engineering notation for Boolean operators when convenient, writing $x_1 y_2 + x_2 z_3$, for example, instead of $(x_1 \wedge y_2) \vee (x_2 \wedge z_3)$.

Quantifying States of Discrete Variables

(Darwiche and Marquis 2021) introduced universal literal quantification for Boolean logic and suggested the following generalization to discrete variables without further study.

Definition 1. For formula Δ and variable X with states x_1, \dots, x_n , the universal quantification of state x_i from Δ is defined as follows: $\forall x_i \cdot \Delta = (\Delta|x_i) \wedge \bigwedge_{j \neq i} (x_j \vee \Delta|x_j)$.

Quantification is commutative so $\forall x_i, y_j \cdot \Delta$, $\forall y_j, x_i \cdot \Delta$ and $\forall \{x_i, y_j\} \cdot \Delta$ are equivalent. We will study and exploit Definition 1 to compute the complete reasons of decisions.

Definition 2. If instance δ is decided positively by class formula Δ , then $\forall \delta \cdot \Delta$ is the ‘complete reason’ for the decision.

The next three results parallel Boolean ones in (Darwiche and Marquis 2021). They are followed by two novel results.

Proposition 1. We have $\forall x_i \cdot \top = \top$ and $\forall x_i \cdot \perp = \perp$; $\forall x_i \cdot x_i = x_i$ and $\forall x_i \cdot \bar{x}_i = \perp$; $\forall x_i \cdot x_j = \perp$ and $\forall x_i \cdot \bar{x}_j = x_i$ when $j \neq i$; $\forall x_i \cdot y_j = y_j$ and $\forall x_i \cdot \bar{y}_j = \bar{y}_j$ when $X \neq Y$.

The next result shows when $\forall x_i$ can be distributed.

Proposition 2. For discrete formulas α, β and state x_i of variable X , we have $\forall x_i \cdot (\alpha \wedge \beta) = (\forall x_i \cdot \alpha) \wedge (\forall x_i \cdot \beta)$. Moreover, if variable X does not occur in both α and β , then $\forall x_i \cdot (\alpha \vee \beta) = (\forall x_i \cdot \alpha) \vee (\forall x_i \cdot \beta)$.

Given Propositions 1 and 2, we can universally quantify states out of \vee -decomposable NNFs in linear time while preserving \vee -decomposability in the resulting NNF.

Proposition 3. Let Δ be a \vee -decomposable NNF and γ be a set of states. Then $\forall \gamma \cdot \Delta$ can be obtained from Δ as follows. For each state $x_i \in \gamma$, replace the occurrences of literals \bar{x}_i , x_j and \bar{x}_j , $j \neq i$, in Δ with \perp , \perp and x_i , respectively.

Consider the class formula $\Delta = \bar{x}_1(x_2 + \bar{y}_1)(\bar{y}_1 + z_1)$ over ternary variables X, Y, Z and instance $\delta = x_2, y_2, z_1$ which is decided positively by Δ . The complete reason for this decision is $\forall \delta \cdot \Delta$. Since Δ is \vee -decomposable, Proposition 3 gives $\forall x_2, y_2, z_1 \cdot \Delta = (x_2)(x_2 + y_2)(y_2 + z_1) = x_2(y_2 + z_1)$. Hence, this instance was decided positively because it has characteristic x_2 and one of two other characteristics, y_2z_1 .

We next identify conditions that allow the distribution of $\forall x_i$ over disjuncts that share variables.

Proposition 4. Consider positive NNFs α, β and state x_i of variable X . If x_i does not occur in α, β , or x_j does not occur in α, β for $j \neq i$, then $\forall x_i \cdot (\alpha \vee \beta) = (\forall x_i \cdot \alpha) \vee (\forall x_i \cdot \beta)$.

For a Boolean variable X with states x and \bar{x} , Proposition 4 says that we can distribute $\forall x$ over disjuncts α and β even if they mention literal \bar{x} (but do not mention x). This is a novel result compared to (Darwiche and Marquis 2021).

Next is another novel condition that licenses the distribution of $\forall x_i$ over disjuncts, which we use to derive closed forms for the complete reasons of decision trees and graphs.

Proposition 5. For NNF α and variable X , we have

$$\forall x_i \cdot (\alpha \vee \bigvee_{x_k \in S} x_k) = (\forall x_i \cdot \alpha) \vee (\forall x_i \cdot \bigvee_{x_k \in S} x_k)$$

if X occurs in α only in disjunctions of the form $\bigvee_{x_k \in S'} x_k$ where $S' \supseteq S$.

Consider variables $X = (x_1, \dots, x_4)$ and $Y = (y_1, y_2)$ and the formulas $\alpha = y_1(x_1 + x_2 + x_4)$ and $\beta = x_1 + x_2$. We can invoke Proposition 5 to distribute $\forall x_1$ using $S = \{x_1, x_2\}$ and $S' = \{x_1, x_2, x_4\}$. Hence, $\forall x_1 \cdot (\alpha \vee \beta) = \forall x_1 \cdot (y_1(x_1 + x_2 + x_4)) \vee \forall x_1 \cdot (x_1 + x_2) = y_1 x_1 + x_1 = x_1$. Propositions 2 and 4 do not license this distribution of $\forall x_i$ though.

The Complete Reasons for Decision Graphs

We next provide closed forms for the complete reasons of decision graphs, which subsume decision trees, in the form of monotone, \vee -decomposable NNFs. This will later facilitate the computation of their prime implicants and implicants (sufficient and necessary reasons). We first treat multi-class decision graphs with nominal features and then treat decision graphs with numeric features; see Figures 1 and 3.

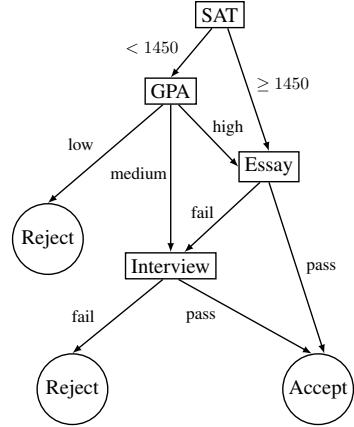


Figure 1: A classifier in the form of a decision graph.

Each leaf node in the graph is labeled with a class c . Each internal node T has edges $\xrightarrow{X, S_1} T_1, \dots, \xrightarrow{X, S_n} T_n$, $n \geq 2$. This node *tests* variable X . Its children are T_1, \dots, T_n and S_1, \dots, S_n is a partition of *some* states of variable X . A decision graph will be represented by its root node. Hence, each node in the graph represents a smaller decision graph. We allow variables to be tested more than once on a path from the root to a leaf but under the following condition, which we call the *weak test-once property*. Consider a path $\dots, T \xrightarrow{X, S_j} T_j, \dots, T' \xrightarrow{X, R_k} T_k, \dots$ in which nodes T and T' test variable X . If no nodes between T and T' on the path test X , then $\{R_k\}_k$ must be a partition of states S_j . Moreover, if T is the first node that tests X on the path, then $\{S_j\}_j$ must be a partition of *all* states for X . For binary variables, the weak test-once property reduces to the standard test-once property: A variable can be tested at most once on any path from the root to a leaf. The weaker property is critical though for treating numeric features as we show later.

A decision graph classifies an instance δ as follows. Suppose $\delta[X]$ is the state of variable X in instance δ . We start at the graph root and repeat the following. When we are at node T with edges $\xrightarrow{X, S_1} T_1, \dots, \xrightarrow{X, S_n} T_n$, we follow the (unique) edge $\xrightarrow{X, S_i} T_i$ which satisfies $\delta[X] \in S_i$. This process leads us to a unique leaf node. The label c of this leaf node is then the class assigned to instance δ .

We next provide a closed-form NNF that captures the instances belonging to some class c in a decision graph.

Definition 3. The NNF for a decision graph T and class c is denoted $\Delta^c[T]$ and defined inductively as follows:

$$\Delta^c[T] = \begin{cases} \top & \text{if } T \text{ has class } c \\ \perp & \text{if } T \text{ has class } c' \neq c \\ \bigwedge_j (\Delta^c[T_j] \vee \bigvee_{x_i \notin S_j} x_i) & \text{if } T \text{ has edges } \xrightarrow{X, S_j} T_j \end{cases}$$

Proposition 6. For decision graph T , class c and instance δ , we have $\delta \models \Delta^c[T]$ iff T assigns class c to instance δ .

This NNF is positive and can be constructed in linear time but is not \vee -decomposable: The disjuncts in $\bigvee_{x_j \notin S_i} x_j$ share variables and variable X will appear in $\Delta^c[T_j]$ if tested again

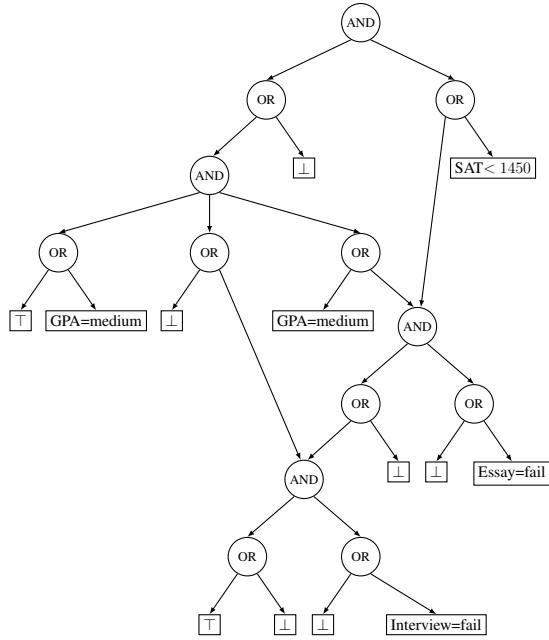


Figure 2: A complete reason constructed by Proposition 7 for the decision graph in Figure 1 and instance $SAT < 1450$, $GPA=\text{medium}$, $\text{Essay}=\text{fail}$, $\text{Interview}=\text{fail}$. This complete reason is in the form of a monotone, \vee -decomposable NNF.

in graph T_j . Yet, this NNF is tractable for universal quantification as revealed in the proof of the next result, which provides closed-form complete reasons for decision graphs.

Proposition 7. Let T be a decision graph, δ be an instance in class c and $\delta[X]$ be the state of variable X in instance δ . The complete reason $\forall \delta \cdot \Delta^c[T]$ is given by the NNF:

$$\Gamma^c[T] = \begin{cases} \top & \text{if } T \text{ has class } c \\ \perp & \text{if } T \text{ has class } c' \neq c \\ \bigwedge_j (\Gamma^c[T_j] \vee \ell_j) & \text{if } T \text{ has edges } \xrightarrow[X, S_j]{} T_j \end{cases} \quad (1)$$

where $\ell_j = \delta[X]$ if $\delta[X] \in S_k$ for some $k \neq j$, else $\ell_j = \perp$.

Consider the decision graph in Figure 1 and an applicant who scored < 1450 on the SAT, had a medium GPA and did not pass their essay or interview. This applicant is rejected by the classifier and the complete reason for the decision, as constructed by Proposition 7, is shown in Figure 2.

Proposition 8. Let T be a decision graph and δ be an instance in class c . The complete reason $\forall \delta \cdot \Delta^c[T]$ in Equation 1 is an NNF that is monotone and \vee -decomposable.

Even though we are working with decision graphs that include discrete variables and multiple classes, we get complete reasons in the form of monotone NNFs, which are effectively Boolean NNFs. This will simplify the computation of necessary and sufficient reasons in later sections, as it allows us to avoid certain complications that can arise when binarizing discrete variables; see, e.g., (Choi et al. 2020).

Numeric Features

Suppose we have a continuous variable X that is being tested at node T in a decision graph. The test will have

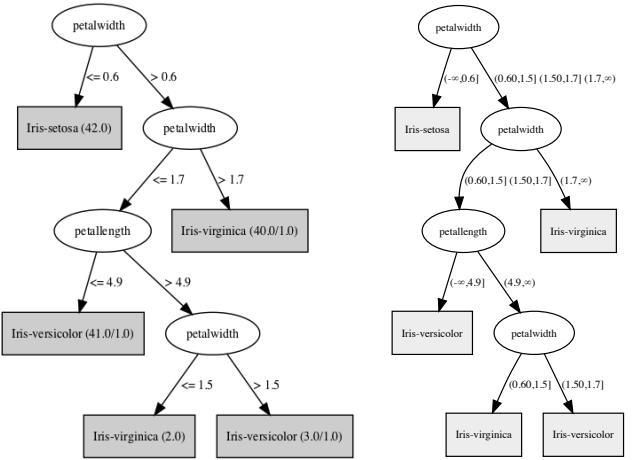


Figure 3: A decision tree with continuous variables learned using weka (left) and its discretization (right).

the form $X \leq t_i$, where t_i is a threshold in $(-\infty, \infty)$. Node T will then have two outgoing edges, one is followed when $X \leq t_i$ (high edge) and the other is followed when $X > t_i$ (low edge); see Figure 3. Suppose now that t_1, \dots, t_n is the set of all thresholds for variable X in the decision graph and assume that these thresholds are in increasing order. We can then treat variable X as a discrete variable with the following $n + 1$ states: $(-\infty, t_1], (t_1, t_2], \dots, (t_{n-1}, t_n], (t_n, \infty)$. If variable X is being tested first at node T , we label the high edge of node T with states $(-\infty, t_1], (t_1, t_2], \dots, (t_{i-1}, t_i]$ and its low edge with states $(t_i, t_{i+1}], \dots, (t_{n-1}, t_n], (t_n, \infty)$. Consider Figure 3 (left). Variable “petalwidth” (W) has three thresholds 0.6, 1.5, 1.7, leading to four discrete states $S_W = (-\infty, 0.6], (0.6, 1.5], (1.5, 1.7], (1.7, \infty)$. Variable “petal-length” (L) has one threshold 4.9, leading to two discrete states $S_L = (-\infty, 4.9], (4.9, \infty)$. Variable W is tested three times in the decision tree. The first test ($W \leq 0.6$) splits states S_W into $S_1 = (-\infty, 0.6]$ for the high edge and $S_2 = (0.6, 1.5], (1.5, 1.7], (1.7, \infty)$ for the low edge. The second test ($W \leq 1.7$) splits S_2 into $S_{21} = (0.6, 1.5], (1.5, 1.7]$ for the high edge and $S_{22} = (1.7, \infty)$ for the low edge. The third and final test ($W \leq 1.5$) splits states S_{21} into $(0.6, 1.5]$ and $(1.5, 1.7]$. The resulting decision tree with discrete variables does not satisfy the test-once property but does satisfy the weak test-once property as shown in Figure 3(right).

Consider now instance $\delta_1 : W = 0.8, L = 5.3$ which is classified as “Iris-virginica” by the decision tree with continuous variables (T_1). We can view this instance as the discrete instance $\delta_2 : W = (0.6, 1.5], L = (4.9, \infty)$ since $0.8 \in (0.6, 1.5]$ and $5.3 \in (4.9, \infty)$. The decision tree with discrete variables (T_2) will also classify instance δ_2 as “Iris-virginica.” A continuous instance and its corresponding discrete instance will be classified identically by decision trees T_1 and T_2 because T_1 cannot discriminate continuous values that belong to the same interval. Finally, to generate the complete reason for instance δ_1 , we compute $\forall \delta_2 \cdot \Delta^c[T_2]$ using Proposition 7 where c is class “Iris-virginica.”

Further Extensions

The closed-form complete reason in Proposition 7 applies directly to *Free Binary Decision Diagrams (FBDDs)* (Gegov and Meinel 1994) and *Ordered Binary Decision Diagrams (OBDDs)* (Bryant 1986) as they are special cases of decision graphs. FBDDs use binary variables and binary classes (\top and \perp). OBDDs are a subset of FBDDs which test variables in the same order along any path from the root to a leaf. We can similarly obtain closed forms for the complete reasons of *Sentential Decision Diagrams (SDDs)* (Darwiche 2011), which test on formulas (sentences) instead of variables. This is possible since given an SDD for Δ we can obtain an SDD for $\neg\Delta$ in linear time. An SDD Δ is an \wedge -decomposable NNF that represents instances for class \top . The SDD for $\neg\Delta$ is also an \wedge -decomposable NNF but represents instances for class \perp . If we negate Δ and $\neg\Delta$ using deMorgan’s law, we obtain \vee -decomposable NNFs for classes \perp and \top , respectively. This allows us to obtain a closed-form, monotone, \vee -decomposable complete reason for any instance using universal quantification. (Darwiche and Marquis 2021) showed that an SDD can be universally quantified in linear time. Earlier, (Darwiche and Hirth 2020) showed that Decision-DNNFs (Huang and Darwiche 2007) can be universally quantified in linear time as well.² Decision-DNNFs cannot be negated efficiently so they do not permit closed-form complete reasons unless we have Decision-DNNFs for classes \top and \perp . While decision tree classifiers are normally learned from data, classifiers such as OBDDs and SDDs are compiled from other classifiers like Bayesian/neural networks and random forests; see, e.g., (Shih, Choi, and Darwiche 2019; Shi et al. 2020; Choi et al. 2020). The relative succinctness of these representations of classifiers has been well studied. FBDDs are a subset of Decision-DNNFs and there is a quasipolynomial simulation of Decision-DNNFs by equivalent FBDDs (Beame et al. 2013). SDDs and FBDDs are not comparable (Beame and Liew 2015; Bollig and Buttkus 2019) so SDDs and Decision-DNNFs are not comparable either. SDDs are exponentially more succinct than OBDDs (Bova 2016).

Necessary and Sufficient Reasons

As mentioned earlier, the prime implicants of a complete reason can be interpreted as *sufficient reasons* for the decision. We next show that the prime implicants of a complete reason can be interpreted as *necessary reasons* for the decision and correspond to contrastive explanations (Ignatiev et al. 2020). We first provide further insights into complete reasons which will help in justifying this interpretation.

Definition 4. Instances δ_1 and δ_2 are ‘congruent’ iff $\delta_1 \cap \delta_2 \models \Delta$ for some class formula Δ . We also say in this case that the decisions on instances δ_1 and δ_2 are congruent.

If instances δ_1 and δ_2 are congruent, they must belong to the same class since $\delta_1 \models \Delta$ and $\delta_2 \models \Delta$ so they are decided

²This was called *consensus* and *filtering* in (Darwiche and Hirth 2020) which is equivalent to universal quantification; see (Darwiche and Marquis 2021). Decision-DNNFs are \wedge -decomposable NNFs in which disjunctions have the form $(x \wedge \alpha) \vee (\bar{x} \wedge \beta)$.

similarly. Moreover, their common characteristics $\delta_1 \cap \delta_2$ are sufficient to justify the decision. That is, the decisions on them are equal and have a common justification.

Proposition 9. Let $\forall\delta \cdot \Delta$ be the complete reason for the decision on instance δ . Then instance δ^* is congruent to instance δ iff $\delta^* \models \forall\delta \cdot \Delta$.

Hence, the complete reason $\forall\delta \cdot \Delta$ captures all, and only, instances that are congruent to instance δ . The complete reason does not capture all instances that are decided similarly to δ since some of these instances may be decided that way for a different reason (the decisions are not congruent).

Consider the class formula $\Delta = \bar{x}_1(x_2 + \bar{y}_1)(\bar{y}_1 + z_1)$ over ternary variables X, Y and Z . The instance $\delta = x_2y_2z_1$ is decided positively by this formula ($\delta \models \Delta$) and the complete reason for this decision is $\forall x_2, y_2, z_1 \cdot \Delta = x_2(y_2 + z_1)$. There are four other instances that satisfy this complete reason, $x_2y_2z_2, x_2y_2z_3, x_2y_1z_1$ and $x_2y_3z_1$. All are decided positively by Δ and the states each share with instance δ justify the decision. Instance $x_3y_2z_1$ is also decided positively by Δ but for a different reason: the states y_2z_1 it shares with instance δ do not justify the decision, $y_2z_1 \not\models \Delta$. Hence, this instance is not captured by the complete reason for δ .

Implicants and Implicates as Reasons

We next review the interpretation of prime implicants as sufficient reasons and discuss the interpretation of prime implicates as necessary reasons for a decision. We will represent these notions by sets of literals, which are interpreted as conjunctions for prime implicants (terms) and as disjunctions for prime implicants (clauses).

Proposition 10. The prime implicants and prime implicants of a complete reason $\forall\delta \cdot \Delta$ are subsets of instance δ .

A prime implicant σ of the complete reason $\forall\delta \cdot \Delta$ can be viewed as a sufficient reason for the underlying decision as it is a minimal subset of instance δ that is guaranteed to sustain the decision, congruently. If we change any part of the instance but for σ , the decision will stick and for a common reason since the new and old instances are congruent. Consider the complete reason in Figure 2 which corresponds to a reject decision on the instance SAT < 1450, GPA=medium, Essay=fail, Interview=fail. There are two prime implicants for this complete reason: (SAT < 1450, GPA=medium, Interview=fail) and (Essay=fail, Interview=fail). Each of these prime implicants is a minimal subset of the instance that is sufficient to sustain the reject decision.

A prime implicate σ of the complete reason $\forall\delta \cdot \Delta$ can be viewed as a necessary reason for the underlying decision as it is a minimal subset of the instance that is essential for sustaining a congruent decision. If we change all states in σ , the decision on the new instance will be different or will be made for a different reason since the new and old instances will not be congruent (we provide a stronger semantics later). Consider again the complete reason in Figure 2 and the corresponding instance and reject decision. There are three prime implicants for this complete reason: (Interview=fail), (SAT < 1450, Essay=fail) and (GPA=medium, Essay=fail). Changing Interview to pass will change the decision. Changing SAT to ≥ 1450 and Essay to pass will

also change the decision. Since GPA is a ternary variable, there are two ways to change its value. If we change GPA and Essay to high and pass, respectively, the decision will change. But if we change these features to low and pass, respectively, the decision will not change but the new instance ($SAT < 1450$, GPA=low, Essay=pass, Interview=fail) will not be congruent with the original instance ($SAT < 1450$, GPA=medium, Essay=fail, Interview=fail). That is, the common characteristics of these instances ($SAT < 1450$, Interview=fail) cannot on their own justify the reject decision.

For yet another example, consider again class formula $\Delta = \bar{x}_1(x_2 + \bar{y}_1)(\bar{y}_1 + z_1)$ over ternary variables X, Y and Z . The complete reason for positive instance $\delta = x_2y_2z_1$ is $\Gamma = \forall \delta \cdot \Delta = x_2(y_2 + z_1)$. The prime implicants of Γ are x_2y_2 and x_2z_1 , which are the sufficient reasons for the decision. If we change instance δ while keeping one of these reasons intact, the decision sticks. The prime implicants of Γ are x_2 and $y_2 + z_1$, which are the necessary reasons for the decision. If we violate one of these reasons, the decision will be different or made for a different reason. Changing instance δ to $x_2y_1z_3$ violates the necessary reason $y_2 + z_1$, which leads to a negative decision. Changing the instance to $\delta^* = x_3y_2z_1$ violates the necessary reason x_2 . The decision remains positive though but for a different reason than why δ is positive. That is, the common characteristics $\delta \cap \delta^* = y_2z_1$ do not justify the decision on these instances, $y_2z_1 \not\models \Delta$.

More on Necessity

We next show that necessary reasons correspond to *basic* contrastive explanations as formalized in (Ignatiev et al. 2020) using the following definition (modulo notation).

Definition 5. Let δ be an instance decided positively by class formula Δ . A ‘contrastive explanation’ of this decision is a minimal subset γ of instance δ such that $\delta \setminus \gamma \not\models \Delta$.

That is, it is possible to change the decision on instance δ by *only* changing the states in γ . Moreover, we must change *all* states in γ for the decision to change.

Proposition 11. Let δ be an instance decided positively by class formula Δ . Then γ is a prime implicate of the complete reason $\forall \delta \cdot \Delta$ iff γ is a contrastive explanation.

This correspondence is perhaps not too surprising given the duality between abductive and contrastive explanations (Ignatiev et al. 2020) and the classical duality between prime implicants and prime implicants. However, it does provide further insights into contrastive explanations: changing the states of a contrastive explanation leads to a non-congruent decision. It also provides further insights on the necessity of prime implicants: while violating a necessary reason will only lead to an instance that is not congruent (decided differently or for a different reason), there must exist at least one violation of each necessary reason which is guaranteed to change the decision. This follows directly from Definition 5. If the variables of a necessary reason are all binary, there is only one way to violate the reason (by negating each variable in the reason). In this case, violating the necessary reason is guaranteed to change the decision.

For an example, let us revisit the complete reason in Figure 2 and the corresponding instance and reject decision.

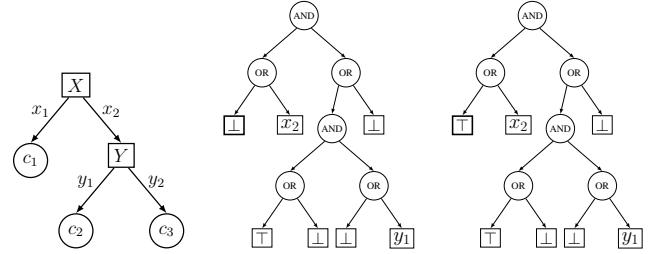


Figure 4: (a) decision tree (b) complete reason for “why x_2y_1 is c_2 ” (c) complete reason for “why x_2y_1 is not c_3 .”

This decision has three necessary reasons: (Interview=fail), ($SAT < 1450$, Essay=fail) and (GPA=medium, Essay=fail). There is only one way to violate each of the first two reasons, and each violation leads to reversing the decision as we saw earlier. There are two ways to violate the third necessary reason. One of these violations (GPA=high, Essay=pass) reverses the decision but the other violation (GPA=low, Essay=pass) keeps the reject decision intact.

In summary, a necessary reason (contrastive explanation) identifies a minimal subset of the instance which is guaranteed to change the decision if that subset is altered *properly*. The minimality condition ensures that we must alter every variable in a necessary reason to change the decision, but it does not specify how to alter it (except for binary variables).

Targeting a Particular Class

Beyond basic contrastive explanations, (Ignatiev et al. 2020) discussed *targeted* contrastive explanations which aim to change the instance class from c to some c^* ; see (Lipton 1990; Miller 2019). This notion is particularly relevant to multi-class classifiers as it corresponds to basic contrastive explanations when the classifier has only two classes. Targeted contrastive explanations can be obtained using the complete reason for *why* the instance was classified as *not* c^* . This can be done using a slight modification of Equation 1 where we modify the first two conditions as follows:

$$\Gamma^c[T] = \begin{cases} \top & \text{if } T \text{ has class } c' \neq c^* \\ \perp & \text{if } T \text{ has class } c^* \\ \bigwedge_j (\Gamma^c[T_j] \vee \ell_j) & \text{if } T \text{ has edges } \xrightarrow{X, S_j} T_j \end{cases} \quad (2)$$

The prime implicants for this complete reason (i.e., necessary reasons) will then identify minimal subsets of the instance that lead to the targeted class c^* , if altered properly.

Consider the decision tree in Figure 4(a) which has two binary features X and Y and three classes c_1, c_2, c_3 . The instance x_2y_1 is classified as c_2 . The complete reason for this decision, as computed by Equation 1, is shown in Figure 4(b) and has two necessary reasons x_2 and y_1 . If we violate the first necessary reason ($x_2 \rightarrow x_1$), the class changes to c_1 . If we violate the second necessary reason ($y_1 \rightarrow y_2$), the class changes to c_3 . Suppose now we wish to change the class c_2 of this instance particularly to c_3 . The complete reason for “why not c_3 ,” as computed by Equation 2, is shown in Figure 4(c) and has only one necessary reason, y_1 . Violating this necessary reason is guaranteed to change the class to c_3 .

Algorithm 1: Shortest Necessary Reasons (SNRs)

Input: monotone and \vee -decomposable NNF Δ with no constants
Output: all shortest implicates of NNF Δ

```

1: function SNR( $\Delta$ )            $\triangleright$  CACHE initialized to NIL
2:   if CACHE( $\Delta$ )  $\neq$  NIL then return CACHE( $\Delta$ )
3:   else if  $\Delta = x_i$  then snr  $\leftarrow \{\{x_i\}\}$ 
4:   else if  $\Delta = \alpha_1 \vee \dots \vee \alpha_n$  then
5:     snr  $\leftarrow$  SNR( $\alpha_1$ )  $\times \dots \times$  SNR( $\alpha_n$ )
6:   else  $\Delta = \alpha_1 \wedge \dots \wedge \alpha_n$ 
7:     snr  $\leftarrow \bigcup_{\text{IML}(\alpha_i) = \text{IML}(\Delta)} \text{SNR}(\alpha_i)$ 
8:   CACHE( $\Delta$ )  $\leftarrow$  snr
9: return snr

```

Algorithm 2: Shortest Sufficient Reasons (SSRs)

Input: monotone and \vee -decomposable NNF Δ with no constants
Output: all shortest implicates of NNF Δ

```

1: function SSR( $\Delta$ )            $\triangleright$  CACHE initialized to NIL
2:    $k \leftarrow 0$ 
3:   repeat
4:     srr  $\leftarrow$  IMP( $\Delta, k, \{\}$ )
5:      $k \leftarrow k + 1$ 
6:   until srr  $\neq \{\}$ 
7: return srr
8: function IMP( $\Delta, k, \sigma$ )       $\triangleright$  computes  $k$ -implicants for  $\Delta|\sigma$ 
9:   if CACHE( $\Delta, k, \sigma$ )  $\neq$  NIL then return CACHE( $\Delta, k, \sigma$ )
10:   $\Sigma \leftarrow \{\}$ 
11:  if  $\Delta = x$  then
12:    if  $x \in \sigma$  then  $\Sigma \leftarrow \{\{\}\}$            $\triangleright \Delta|\sigma$  valid
13:    else if  $k \geq 1$  then  $\Sigma \leftarrow \{\{x\}\}$ 
14:  else if  $\Delta = \alpha_1 \vee \dots \vee \alpha_n$  then
15:    for  $\Sigma_i \leftarrow$  IMP( $\alpha_i, k, \sigma$ ) do
16:      if  $\Sigma_i \neq \{\{\}\}$  then  $\Sigma \leftarrow \Sigma \cup \Sigma_i$   $\triangleright \alpha_i|\sigma$  not valid
17:      else  $\Sigma \leftarrow \{\{\}\}; \text{break}$            $\triangleright \Delta|\sigma$  valid
18:  else  $\Delta = \alpha_1 \wedge \dots \wedge \alpha_n$ 
19:    for  $\sigma_1 \in$  IMP( $\alpha_1, k, \sigma$ ) do
20:      for  $\sigma_2 \in$  IMP( $\alpha_2 \wedge \dots \wedge \alpha_n, k - |\sigma_1|, \sigma \cup \sigma_1$ ) do
21:         $\Sigma \leftarrow \text{REMOVE\_SUBSUMED}(\Sigma \cup \{\sigma_1 \cup \sigma_2\})$ 
22:  CACHE( $\Delta, k, \sigma$ )  $\leftarrow \Sigma$ 
23: return  $\Sigma$ 

```

Computing Shortest Explanations

A complete reason may have too many prime implicants and implicants. We will therefore provide algorithms for computing the *shortest* implicants and implicants (which must be prime) for monotone, \vee -decomposable NNFs. As discussed earlier, we can in linear time obtain complete reasons in this form for decision graphs and SDDs, which include FBDDs, OBDDs and decision trees as special cases.

Shortest Necessary Reasons. This will be an output polynomial algorithm that is based on three (conceptual) passes on the complete reason which we describe next.

Definition 6. The ‘implicate minimum length (IML)’ of a valid formula is ∞ . For non-valid formulas, it is the minimum length attained by any implicate of the formula.

The first pass computes the implicate minimum length.

Proposition 12. The IML of a monotone, \vee -decomposable NNF is computed as follows: $\text{IML}(\top) = \infty$, $\text{IML}(\perp) = 0$,

$$\text{IML}(x_i) = 1, \text{IML}(\alpha \vee \beta) = \text{IML}(\alpha) + \text{IML}(\beta) \text{ and } \text{IML}(\alpha \wedge \beta) = \min(\text{IML}(\alpha), \text{IML}(\beta)).$$

The second pass prunes the NNF using the IML of nodes.

Proposition 13. Let PRUNE(Δ) be the NNF obtained from monotone, \vee -decomposable NNF Δ by dropping α_i from conjunctions $\alpha = \alpha_1 \wedge \dots \wedge \alpha_n$ if $\text{IML}(\alpha_i) > \text{IML}(\alpha)$. Then PRUNE(Δ) is a monotone, \vee -decomposable NNF and its prime implicants are the shortest implicants of Δ .

The third pass computes the prime implicants of NNF PRUNE(Δ) in output polynomial time. Algorithm 1 implements the second and third passes assuming the first, linear-time pass has been performed. It represents an implicate by a set of literals and uses the Cartesian product operation on sets of implicants: $S_1 \times S_2 = \{\sigma_1 \cup \sigma_2 \mid \sigma_1 \in S_1, \sigma_2 \in S_2\}$. Algorithm 1 applies the second pass implicitly by excluding conjuncts on Line 7. This is the standard procedure for computing the prime implicants of a monotone NNF, but with no subsumption checking which is critical for its complexity. In the standard procedure, one must ensure that the implicants computed on Lines 5 and 7 are reduced: no implicate σ_1 subsumes another σ_2 ($\sigma_1 \subseteq \sigma_2$).³ Since NNF PRUNE(Δ) is \vee -decomposable, the disjuncts $\alpha_1, \dots, \alpha_n$ on Line 5 do not share variables. Hence, if every $\text{SNR}(\alpha_i)$ is reduced, their Cartesian product is reduced. Moreover, due to pruning in the second pass, the implicants $\text{SNR}(\alpha_i)$ computed on Line 7 all have the same length so no subsumption is possible.

Proposition 14. Let Δ be a monotone, \vee -decomposable NNF with M shortest implicants, N nodes and E edges. The time complexity of SNR(Δ) in Algorithm 1 is $O(M \cdot E)$ and its space complexity is $O(M \cdot N)$.

We obtain a tighter complexity if we apply Algorithm 1 to the closed-form complete reasons of decision trees given by Proposition 7, due to the following bound on the number of prime implicants (a superset of shortest implicants).

Proposition 15. For a decision tree, the complete reason for an instance in class c has $\leq L$ prime implicants, where L is the number of leaves in the tree labeled with a class $c' \neq c$.

The complete reason for a decision tree T has $O(|T|)$ nodes and edges, where $|T|$ is the decision tree size (see Proposition 7). The time and space complexity of Algorithm 1 is then $O(|T| \cdot L)$ for decision trees.

(Huang et al. 2021b) showed that the number of contrastive explanations is linear in the decision tree size. Proposition 15 tightens this result by providing a more specific bound. For a decision T with binary variables and binary classes, (Audemard et al. 2021) showed that the set of all contrastive explanations can be computed in time polynomial in $|T| + n$, where n is the number of variables. Algorithm 1 comes with a tighter complexity for the computation of shortest contrastive explanations for decision trees and applies to multi-class decision trees with discrete features.

³One can compute the prime implicants of a monotone NNF by simply converting it to a CNF and then removing subsumed clauses. Similarly, one can compute the prime implicants of a monotone NNF by converting it to a DNF and removing subsumed terms. See, e.g., (Crama and Hammer 2011).

benchmark		decision tree/graph properties						SR		SSR		NR		SNR	
name	examples	nodes	num	nom	classes	card	acc	count	time	count	time	count	time	count	time
adult	48842	726	6	7	2	24	86.0	2.8	0.0005	1.4	0.0007	5.6	0.0004	2.9	0.0003
compas	5278	55	5	3	2	8	71.2	1.7	0.0002	1.1	0.0003	2.9	0.0002	1.8	0.0001
fash-mnist	70000	6681	734	0	10	29	80.8	1851.6	3.7925 ⁶⁵⁷	7.5	0.0348 ⁸	104.8	0.0123	12.4	0.0008
gisette	7000	231	111	0	2	3	93.9	3288.4	6.2361 ⁴⁸⁶	95.5	0.0545	31.8	0.0013	7.4	0.0004
isolet	7797	645	201	0	26	7	83.9	8.9	0.0008	1.5	0.0026	17.7	0.0006	10.1	0.0003
la1s.wc	3204	457	201	0	6	3	73.2	461.1	1.8227 ³⁶	7.7	0.0368	42.1	0.0038	19.3	0.0004
mnist-784	70000	4365	477	0	10	18	88.3	1833.3	4.6191 ⁶²²	5.5	0.0235 ¹	103.9	0.0099	11.5	0.0008
nomao	34465	932	61	25	2	17	95.3	1640.8	4.9576 ¹⁵⁶	2.4	0.0180	38.6	0.0018	4.3	0.0005
ohscal.wc	11162	1761	582	0	10	6	70.9	787.1	2.0714 ²²⁵	86.2	0.3018 ²¹	59.9	0.0090	23.7	0.0006
spambase	4601	189	37	0	2	8	91.5	35.8	0.0029	4.6	0.0060	16.9	0.0007	4.1	0.0003
andes	—	5454	—	24	2	2	—	78.8	0.1660	2.6	0.0295	58.0	0.0300	4.5	0.0055
emdec6g30	—	4154	—	30	2	2	—	11.6	0.0362	1.8	0.0445	13.6	0.0184	2.9	0.0048
math-skills	—	3693629	—	46	2	2	—	9.6	0.2895 ²⁶	7.6	0.0261	9.9	0.4271 ¹⁸	3.4	0.0290
mooring	—	14468	—	22	2	2	—	145.4	2.3991 ¹⁴	20.9	3.9951 ⁷	216.3	0.7678 ¹	16.2	0.0189
tcc4e38	—	22508	—	38	2	2	—	14.4	0.1370	2.6	0.0155	10.4	0.0632	2.0	0.0103

Table 1: Evaluating Algorithms 1 & 2. Times in secs. First ten entries are decision trees. Last five entries are decision graphs.

We finally observe that if an NNF is monotone and \wedge -decomposable, then one can develop a dual of Algorithm 1 for computing the shortest prime implicants of the NNF.

Shortest Sufficient Reasons We next present Algorithm 2 for computing the shortest implicants of monotone, \vee -decomposable NNFs which is a hard task. For decision trees, the problem of deciding whether there exists a sufficient reason of length $\leq k$ is NP-complete (Barceló et al. 2020). Since decision trees have closed-form complete reasons that are monotone and \vee -decomposable, computing the shortest implicants for this class of NNFs is hard. (Audevard et al. 2021) showed that the number of shortest sufficient reasons for decision trees can be exponential and provided an incremental algorithm for computing the shortest sufficient reasons for decision trees with binary variables and binary classes, based on a reduction to the PARTIAL MAXSAT problem. Algorithm 2 has a broader scope, does not require a reduction and is based on two key techniques.

The first technique is to compute all unsubsumed implicants of length $\leq k$, called k -implicants, starting with $k = 0$. If no implicants are found, k is incremented and the process is repeated. The second technique relates to computing the k -implicants of a conjunction $\alpha \wedge \beta$. If we have the k -implicants S for α and the k -implicants R for β , we can compute the Cartesian product $S \times R$ and keep unsubsumed implicants of length $\leq k$. Algorithm 2 does something more refined. It first computes the k -implicants S for α . For each implicant $\sigma \in S$, it then computes and accumulates the k' -implicants for $\beta|\sigma$ where $k' = k - |\sigma|$. These techniques control the number of generated k -implicants at each NNF node (smaller k leads to fewer k -implicants). Our implemented caching scheme on Lines 9 & 22 exploits the following properties. If the k -implicants for $\Delta|\sigma$ are $\{\{\}\}$, then these are also its j -implicants for all j . Further, if we cached the k -implicants for $\Delta|\sigma$, then we can use them to retrieve its j -implicants for any $j \leq k$ by selecting implicants of length $\leq j$. We empirically evaluate Algorithms 1 & 2 next.

Empirical Evaluation. Table 1 depicts an empirical eval-

uation on decision trees learned from OpenML datasets (Vanschoren et al. 2013) and binary decision graphs compiled from Bayesian network classifiers (Shih, Choi, and Darwiche 2018a, 2019). The decision trees were learned by WEKA (Frank et al. 2010) using python-weka-wrapper3 available at pypi.org. We used WEKA’s J48 classifier with default settings, which learns pruned C4.5 decision trees with numeric and nominal features (Quinlan 1993). Each dataset was split using WEKA into training (85%) and testing (15%) data. We aimed for OpenML datasets with more than 100 features since many smaller datasets we tried were very easy, but we kept a few smaller ones since they are commonly reported on (adult, compas, spambase). Some of the learned decision trees had significantly fewer variables than the corresponding datasets (e.g., gisette has 5000 features but the learned decision tree has 111). The decision graphs we used are the reportedly largest ones compiled by (Shih, Choi, and Darwiche 2018a, 2019). For each decision tree, we computed reasons for decisions on 1000 instances sampled from testing data (or all testing data if smaller than 1000). We tried random instances but they were much easier. For each decision graph, we computed complete reasons for 1000 random instances (there is no corresponding data). The total number of instances for the fifteen benchmarks was 13963. We did not report the time for computing a complete reason as this is a closed form with linear size (Equation 1).

We compared four algorithms: SNR (Algorithm 1), NR (standard algorithm for computing prime implicants of a monotone NNF but with no subsumption checking at \vee -nodes since the input NNF is \vee -decomposable),⁴ SSR (Algorithm 2) and SR (dual of NR). Each instance had a timeout of 60 seconds. In Table 1, nodes, num, nom, classes; card and acc stand for number of nodes, numeric features, nominal features, classes; maximum cardinality of variables and ac-

⁴More precisely, NR is Algorithm 1 with two exceptions. First, the NNF is not pruned on Line 7 so the union is over all α_i . Second, subsumption checking is applied after Line 7 to ensure that all computed implicants are subset-minimal.

curacy. Count and time are averages over instances that both SR/SSR (NR/SNR) finished. The bolded exponent of time is the number of instances that timed out (not reported if zero). The supplementary material contains further statistics such as stdev, mean and max. We used a Python implementation on a dual Intel(R) Xeon E5-2670 CPUs running at 2.60GHz and 256GB RAM. As revealed by Table 1, SSR is quite effective. Its average running time is normally in milliseconds, it timed out on only 37 instances and can be two orders of magnitude faster than SR which timed out on 2222 instances. SNR is also much faster than NR but the latter is also very effective on decision trees (see Proposition 15) but timed out on 19 decision graph instances. All algorithms are quite effective on the easier benchmarks.

Conclusion

We studied the computation of complete reasons for multi-class classifiers with nominal and numeric features. We derived closed forms for the complete reasons of decision trees and graphs in the form of monotone, \vee -decomposable NNFs and showed how similar forms can be derived for SDDs. We further established a correspondence between the prime implicants of complete reasons and contrastive explanations. We then presented an output polynomial algorithm for enumerating the shortest implicants (shortest necessary reasons) for complete reasons in the above form. We also presented a simple algorithm for enumerating the shortest implicants (shortest sufficient reasons) which appears to be effective based on an empirical evaluation over fifteen datasets.

Acknowledgements

This work has been partially supported by NSF grant #ISS-1910317 and ONR grant #N00014-18-1-2561.

References

- Audemard, G.; Bellart, S.; Bounia, L.; Koriche, F.; Lagniez, J.-M.; and Marquis, P. 2021. On the Explanatory Power of Decision Trees. *CoRR*, <https://arxiv.org/pdf/2108.05266.pdf>.
- Audemard, G.; Koriche, F.; and Marquis, P. 2020. On Tractable XAI Queries based on Compiled Representations. In *KR*, 838–849.
- Barceló, P.; Monet, M.; Pérez, J.; and Subercaseaux, B. 2020. Model Interpretability through the lens of Computational Complexity. In *NeurIPS*.
- Beame, P.; Li, J.; Roy, S.; and Suciu, D. 2013. Lower Bounds for Exact Model Counting and Applications in Probabilistic Databases. In *UAI*. AUAI Press.
- Beame, P.; and Liew, V. 2015. New Limits for Knowledge Compilation and Applications to Exact Model Counting. In *UAI*, 131–140. AUAI Press.
- Bollig, B.; and Buttkus, M. 2019. On the Relative Succinctness of Sentential Decision Diagrams. *Theory Comput. Syst.*, 63(6): 1250–1277.
- Bova, S. 2016. SDDs Are Exponentially More Succinct than OBDDs. In *AAAI*, 929–935. AAAI Press.
- Bryant, R. E. 1986. Graph-Based Algorithms for Boolean Function Manipulation. *IEEE Trans. Computers*, 35(8): 677–691.
- Chan, H.; and Darwiche, A. 2003. Reasoning about Bayesian Network Classifiers. In *UAI*, 107–115. Morgan Kaufmann.
- Choi, A.; Shih, A.; Goyanka, A.; and Darwiche, A. 2020. On Symbolically Encoding the Behavior of Random Forests. *CoRR*, abs/2007.01493.
- Choi, A.; Xue, Y.; and Darwiche, A. 2012. Same-decision probability: A confidence measure for threshold-based decisions. *Int. J. Approx. Reason.*, 53(9): 1415–1428.
- Crama, Y.; and Hammer, P. L. 2011. *Boolean Functions - Theory, Algorithms, and Applications*, volume 142 of *Encyclopedia of mathematics and its applications*. Cambridge University Press.
- Darwiche, A. 2011. SDD: A New Canonical Representation of Propositional Knowledge Bases. In *IJCAI*, 819–826. IJCAI/AAAI.
- Darwiche, A.; and Hirth, A. 2020. On the Reasons Behind Decisions. In *ECAI*, volume 325 of *Frontiers in Artificial Intelligence and Applications*, 712–720. IOS Press.
- Darwiche, A.; and Marquis, P. 2021. On Quantifying Literals in Boolean Logic and Its Applications to Explainable AI. *J. Artif. Intell. Res.*, 72: 285–328.
- Frank, E.; Hall, M. A.; Holmes, G.; Kirkby, R.; Pfahringer, B.; Witten, I. H.; and Trigg, L. 2010. Weka-A Machine Learning Workbench for Data Mining. In *Data Mining and Knowledge Discovery Handbook*, 1269–1277. Springer.
- Gergov, J.; and Meinel, C. 1994. Efficient Boolean Manipulation With OBDD’s can be Extended to FBDD’s. *IEEE Trans. Computers*, 43(10): 1197–1209.
- Huang, J.; and Darwiche, A. 2007. The Language of Search. *J. Artif. Intell. Res.*, 29: 191–219.
- Huang, X.; Izza, Y.; Ignatiev, A.; Cooper, M. C.; Asher, N.; and Marques-Silva, J. 2021a. Efficient Explanations for Knowledge Compilation Languages. *CoRR*, abs/2107.01654.
- Huang, X.; Izza, Y.; Ignatiev, A.; and Marques-Silva, J. 2021b. On Efficiently Explaining Graph-Based Classifiers. *CoRR*, abs/2106.01350.
- Ignatiev, A.; Narodytska, N.; Asher, N.; and Marques-Silva, J. 2020. From Contrastive to Abductive Explanations and Back Again. In *AI*IA*, volume 12414 of *Lecture Notes in Computer Science*, 335–355. Springer.
- Ignatiev, A.; Narodytska, N.; and Marques-Silva, J. 2019a. Abduction-Based Explanations for Machine Learning Models. In *Proceedings of the Thirty-Third Conference on Artificial Intelligence (AAAI)*, 1511–1519.
- Ignatiev, A.; Narodytska, N.; and Marques-Silva, J. 2019b. On Validating, Repairing and Refining Heuristic ML Explanations. *CoRR*, abs/1907.02509.
- Lipton, P. 1990. Contrastive Explanation. *Royal Institute of Philosophy Supplement*, 27: 247–266.
- Miller, T. 2019. Explanation in artificial intelligence: Insights from the social sciences. *Artif. Intell.*, 267: 1–38.

- Narodytska, N.; Kasiviswanathan, S. P.; Ryzhyk, L.; Sagiv, M.; and Walsh, T. 2018. Verifying Properties of Binarized Deep Neural Networks. In *Proc. of AAAI'18*, 6615–6624.
- Quinlan, R. 1993. *C4.5: Programs for Machine Learning*. San Mateo, CA: Morgan Kaufmann Publishers.
- Ribeiro, M. T.; Singh, S.; and Guestrin, C. 2016. "Why Should I Trust You?": Explaining the Predictions of Any Classifier. In *KDD*, 1135–1144. ACM.
- Ribeiro, M. T.; Singh, S.; and Guestrin, C. 2018. Anchors: High-Precision Model-Agnostic Explanations. In *AAAI*, 1527–1535. AAAI Press.
- Shi, W.; Shih, A.; Darwiche, A.; and Choi, A. 2020. On Tractable Representations of Binary Neural Networks. In *KR*, 882–892.
- Shih, A.; Choi, A.; and Darwiche, A. 2018a. Formal Verification of Bayesian Network Classifiers. In *PGM*, volume 72 of *Proceedings of Machine Learning Research*, 427–438. PMLR.
- Shih, A.; Choi, A.; and Darwiche, A. 2018b. A Symbolic Approach to Explaining Bayesian Network Classifiers. In *IJCAI*, 5103–5111. ijcai.org.
- Shih, A.; Choi, A.; and Darwiche, A. 2019. Compiling Bayesian Network Classifiers into Decision Graphs. In *AAAI*, 7966–7974. AAAI Press.
- Vanschoren, J.; van Rijn, J. N.; Bischl, B.; and Torgo, L. 2013. OpenML: Networked Science in Machine Learning. *SIGKDD Explorations*, 15(2): 49–60.
- Wang, E.; Khosravi, P.; and den Broeck, G. V. 2021. Probabilistic Sufficient Explanations. In *IJCAI*, 3082–3088. ijcai.org.