

Sufficient Reasons for Classifier Decisions in the Presence of Domain Constraints

Niku Gorji, Sasha Rubin

School of Computer Science, The University of Sydney, Australia
niku.gorji@sydney.edu.au, sasha.rubin@sydney.edu.au

Abstract

Recent work has unveiled a theory for reasoning about the decisions made by binary classifiers: a classifier describes a Boolean function, and the reasons behind an instance being classified as positive are the prime-implicants of the function that are satisfied by the instance. One drawback of these works is that they do not explicitly treat scenarios where the underlying data is known to be constrained, e.g., certain combinations of features may not exist, may not be observable, or may be required to be disregarded. We propose a more general theory, also based on prime-implicants, tailored to taking constraints into account. The main idea is to view classifiers as describing partial Boolean functions that are undefined on instances that do not satisfy the constraints. We prove that this simple idea results in more parsimonious reasons. That is, not taking constraints into account (e.g., ignoring, or taking them as negative instances) results in reasons that are subsumed by reasons that do take constraints into account. We illustrate this improved succinctness on synthetic classifiers and classifiers learnt from real data.

1 Introduction

A recent line of enquiry for providing reasons for classifier decisions is on supplying principled reasons for individual instances with formal guarantees of subset- or cardinality-minimality (Shih, Choi, and Darwiche 2018; Darwiche and Hirth 2020; Ignatiev, Narodytska, and Marques-Silva 2019). Contrary to the more scalable heuristic approaches (Ribeiro, Singh, and Guestrin 2016; Lakkaraju et al. 2019; Iyer et al. 2018), these principled approaches guarantee the quality of produced reasons and therefore can serve to validate, benchmark, and potentially provide insights for improving the solutions of the heuristic approaches.

A noteworthy drawback of these methods however, is that they do not deal with reasoning in the presence of constraints or background knowledge.

Constraints may arise from the structure and interdependencies between features present in data (Darwiche 2020). As a simple example, consider a medical setting in which some combinations of drugs are never prescribed together and thus will not appear in any dataset: if we know that drug A and drug B are never prescribed together (i.e.,

the constraint), then a reason of the form “drug A was prescribed and drug B was not prescribed” is overly redundant; however, the current methods in the state of the art produce exactly such overly complicated reasons. We argue that in this situation, it is more parsimonious to supply the reason “drug A was prescribed”. Although it might be obvious how to process reasons to take such simple constraints into account, it is by no means obvious how to handle semantic constraints represented by arbitrarily complex Boolean formulas. It is the purpose of this work to provide a simple and elegant approach for doing this. More precisely, we address the problem of how to incorporate background knowledge, specifically input (domain) constraints, into supplying reasons behind individual decisions of classifiers.

Our contributions are as follows:

1. We provide a crisp formalisation of reasons for classifier decisions that takes constraints into account, resulting in reasons that are at least (and sometimes more) parsimonious, i.e., more general and more succinct, than not taking constraints into account. The central insight, both simple and powerful, is to treat a classifier as a partial function that is not defined on input instances that do not satisfy the constraint, and then to use the classic definition of prime-implicant on partial functions (Coudert 1994) as the instantiation of the word “reason”. This immediately and naturally generalises the state of the art from the unconstrained setting to the constrained setting.
2. We provide a simple reduction of the computational problem of finding all reasons of a classifier’s decision for a given instance in the presence of constraints to the same problem in the unconstrained setting. This allows one to re-use existing algorithms and tools from the unconstrained setting. The idea is that if the constraint is given by the Boolean formula κ , and the decision-function by φ , then reasons of φ -decisions that take κ into account are exactly the reasons of $(\kappa \rightarrow \varphi)$ -decisions.¹ We prove that all other variations, including the seemingly natural variation $(\kappa \wedge \varphi)$, provide no more, and sometimes less, parsimonious reasons.
3. We show, both theoretically and empirically on syn-

¹This simple reduction can be done in linear-time for formulas represented as strings or parse-trees, and in polynomial-time for formulas represented as OBDD_< (Darwiche and Marquis 2002).




Figure 1: Workflow

thetic classifiers and classifiers learnt from data, that approaches that ignore constraints may supply reasons that are unnecessarily long since they redundantly encode knowledge already described in the constraints.

The workflow is illustrated in Figure 1.

2 Preliminaries

We begin by recalling just enough logical background to be able to explain our theory.

Boolean logic. Let $\mathbf{X} = \{X_1, X_2, \dots, X_n\}$ be a set of n Boolean variables (aka *features*). The set of *Boolean formulas* is generated from \mathbf{X} , the constants \top (true) and \perp (false), and the logical operations \wedge (conjunction), \vee (disjunction), \neg (negation), \rightarrow (conditional) and \leftrightarrow (bi-conditional). Variables X and their negations $\neg X$ are called *literals*. A *term* t is a conjunction of literals; the empty-conjunction is also denoted \top . An *instance* (over \mathbf{X}) is an element of $\{0, 1\}^n$, and is denoted \mathbf{x} (intuitively, it is an instantiation of the variables \mathbf{X}). An instance \mathbf{x} *satisfies* a formula φ if φ evaluates to true when the variables in φ are assigned truth-values according to \mathbf{x} . The set of instances that satisfy the formula φ is denoted $[\varphi]$, and is called the set *represented* by φ , i.e., a set C of instances is *represented* by φ if $C = [\varphi]$. If $[\varphi] = [\psi]$ then we say that φ, ψ are *logically equivalent*, i.e., they mean the same thing. For terms s, t , we say that s *subsumes* t if $[t] \subseteq [s]$, i.e., if every instance that satisfies t also satisfies s . If $[t] \subset [s]$ then we say that s *properly subsumes* t .

Partial Boolean functions, and prime-implicants. A *partial Boolean function* f (over \mathbf{X}) is a function $\{0, 1\}^n \rightarrow \{0, 1, *\}$. For $i \in \{0, 1, *\}$ define f^i to be the set $f^{-1}(i)$. The instances in f^1 are called the *positive instances* of f . If the set f^* is empty, then f is a *total Boolean function*. For readability, we may denote total Boolean functions by capital letters, typically F . If $[\varphi] = F^1$ we say that the formula φ *represents* the total Boolean function F . A term t is an *implicant* of f if $[t] \subseteq f^1 \cup f^*$; it is *prime* if no other implicant of f subsumes t . Intuitively, t is prime if removing any literal from t results in a term that is no longer an implicant. This generalises the notion of implicant and prime-implicant from total Boolean functions, cf. (Quine 1952; Shih, Choi, and Darwiche 2018; Darwiche and Hirth 2020), to partial Boolean functions, cf. (McCluskey 1956; Coudert 1994).

Decision-functions. Total Boolean functions naturally arise as the decision-functions of threshold-based binary classifiers (Choi et al. 2019; Shih, Choi, and Darwiche

2018): the *decision-function* F of a threshold-based classifier is the function that maps an instance \mathbf{x} to 1 if $Pr(d = 1|\mathbf{x}) \geq T$, and to 0 otherwise; here d is a binary class variable, and Pr is the distribution specified by the classifier, and T is a user-defined classification threshold.

Definition 1 (Sufficient reasons for total functions). (Shih, Choi, and Darwiche 2018; Darwiche and Hirth 2020) Let F be a total Boolean function and let \mathbf{x} be a positive instance of F . A term t is a sufficient reason of the decision $F(\mathbf{x}) = 1$ if (i) t is a prime-implicant of F , and (ii) t is satisfied by \mathbf{x} .

Sufficient reasons were called *PI-explanation* in (Shih, Choi, and Darwiche 2018, Definition 5).

Standard convention. We freely interchange total Boolean functions and the formulas that represent them. In particular, if φ represents the total Boolean function F , we may refer to implicants, prime-implicants, and sufficient reasons of φ (instead of F).

3 Problem Setting

The problem we address is how to define reasons behind decisions of a classifier in the presence of domain constraints.

Definition 2. A constraint is a set C of instances over \mathbf{X} .

We typically represent constraints by Boolean formulas. Here are just a few examples that show that constraints are ubiquitous. In a medical setting, constraints of the form $(X_1 \rightarrow X_2)$ may reflect that people with condition X_1 also have condition X_2 , e.g., X_1 may mean “is pregnant” and X_2 may mean “is a woman”. In a university degree structure: constraints of the form $X_1 \rightarrow (X_2 \wedge X_3)$ may reflect that X_2 and X_3 are prerequisites to X_1 ; constraints of the form $X_1 \rightarrow \neg(X_2 \vee X_3)$ may reflect prohibitions; and constraints of the form $X_1 \wedge X_2$ may reflect compulsory courses. In configuration problems, e.g., that arise when users purchase products, the user may be required to configure their product subject to certain constraints, and constraints of the form $(X_1 \vee X_2) \wedge \neg(X_1 \wedge X_2)$ may reflect that the user needs to choose between two basic models. These constraints also result from one-hot encodings of a categorical variables, e.g., if M is a 12-valued variable representing months, and X_i for $i = 1, \dots, 12$ is Boolean variable, then the induced constraint is $(\bigvee_i X_i) \wedge (\bigwedge_{i \neq j} \neg(X_i \wedge X_j))$. Combinatorial objects have natural constraints, e.g., rankings are ordered sets, trees are acyclic graphs, and games have rules, see Section 6. Finally, the assumption in this paper is that constraints are *hard*, i.e., instances that are not in C will not appear in any data and can be ignored (e.g., they will not appear in training or testing data).

Recall that threshold-based classifiers produce representations of *total* Boolean functions. This suggests the following useful terminology:

Definition 3. A constrained decision-function is a pair (F, C) consisting of a total Boolean function F and a constraint C .

We thus ask:

How should one define reasons behind decisions of constrained decision-functions?

We posit that a suitable notion of “reason” that takes constraints into account satisfies the following desired properties:

- D1. it does not depend on the representations of F or C , i.e., it is a semantic notion;
- D2. it does not depend on the values $F(\mathbf{x})$ for $\mathbf{x} \notin C$, i.e., if F, G agree on C (and perhaps disagree on the complement of C), then reasons for (F, C) should be the same as reasons for (G, C) ;
- D3. in case there are no constraints, i.e., $C = \{0, 1\}^n$, we recover the notion of sufficient reasons from (Darwiche and Hirth 2020; Shih, Choi, and Darwiche 2018), see Definition 1.
- D4. it is not less succinct (and is sometimes more succinct) than not taking constraints into account. Said differently, reasons should not contain redundancies that are captured by the constraint.

We offer a formalisation that satisfies these desiderata.

4 Reasons in the Presence of Constraints

In this section we provide the main definition of reasons in the presence of constraints (Definition 5) and show that it satisfies all the desired properties D1-D4 listed in Section 3.

Desiderata D1 and D2 motivate the fundamental insight that constrained decision-functions should be treated as partial Boolean functions:

Definition 4. For a constrained decision-function (F, C) , let f_C be the partial Boolean function that maps \mathbf{x} to $F(\mathbf{x})$ if $\mathbf{x} \in C$, and to $*$ otherwise.

D1 and D2 imply that the notion of reason should only depend on the partial function f_C . So, we now define sufficient reasons (of decisions) for partial Boolean functions.

Definition 5 (Sufficient reasons for partial functions). Let f be a partial Boolean function and let \mathbf{x} be a positive instance of f . A term t is a sufficient reason of the decision $f(\mathbf{x}) = 1$ if (i) t is a prime-implicant of f , and (ii) t is satisfied by \mathbf{x} .

If f happens to be a total function, then a term t is a sufficient reason for $f(\mathbf{x}) = 1$ according to Definition 5 if and only if it is a sufficient reason for $f(\mathbf{x}) = 1$ according to Definition 1. Thus we see that D3 holds since $C = \{0, 1\}^n$ means that f is a total function.

Remark 1. The difference between sufficient reasons for partial vs total functions is that a term t is an implicant of a partial function f if $[t] \subseteq f^1 \cup f^*$, and not just if $[t] \subseteq f^1$ (as is the case for total functions).

Example 1. Consider the total Boolean function F over $X = \{X_1, X_2\}$ represented by the formula $(X_1 \leftrightarrow X_2)$. Suppose a constraint is specified by the formula $(X_1 \rightarrow X_2)$, thus $C = \{(0, 0), (0, 1), (1, 1)\}$. Table 1 provides both F and the partial Boolean function f_C . The prime-implicants of F are $(X_1 \wedge X_2)$ and $(\neg X_1 \wedge \neg X_2)$. The only sufficient reason of the decision $F(0, 0) = 1$ is the term $(\neg X_1 \wedge \neg X_2)$, and the only sufficient reason of the decision $F(1, 1) = 1$ is the term $(X_1 \wedge X_2)$. The prime-implicants of

X_1	X_2	F	f_C
0	0	1	1
0	1	0	0
1	0	0	*
1	1	1	1

Table 1: Partial Boolean function determined by the total function F and the constraint C specified by $(X_1 \rightarrow X_2)$. The row corresponding to the instance not in the constraint is greyed-out.

f_C are $\neg X_2$ and X_1 . The only sufficient reason of the decision $f_C(0, 0) = 1$ is $\neg X_2$, and the only sufficient reason of the decision $f_C(1, 1) = 1$ is X_1 .

Finally, recall that D4 says that reasons that take constraints into account should not be less succinct than those that do. We formalise this by proving that f_C supplies reasons that are as least as succinct as reasons using F .²

Lemma 1. If f, g are partial functions such that $f^1 \cup f^* \subseteq g^1 \cup g^*$, then every sufficient reason for $f(\mathbf{x}) = 1$ is subsumed by some sufficient reason of $g(\mathbf{x}) = 1$.

Proof. A sufficient reason of $f(\mathbf{x}) = 1$, being an implicant of f , is also an implicant of g (since $f^1 \cup f^* \subseteq g^1 \cup g^*$). Now apply the simple fact that every implicant of a function is subsumed by some prime-implicant of that function. \square

We immediately get the following fundamental comparison of sufficient reasons for constrained vs unconstrained decision-functions:

Theorem 1. Let (F, C) be a constrained decision-function, and suppose \mathbf{x} is a positive instance of f_C . Then every sufficient reason of $F(\mathbf{x}) = 1$ is subsumed by some sufficient reason of $f_C(\mathbf{x}) = 1$.

Proof. Given F and C , apply Lemma 1 taking $f = F$ and $g = f_C$. To see that the hypothesis holds, simply note that $F^1 \cup F^* = F^1$ and $(f_C)^1 \cup (f_C)^* = (F^1 \cap C) \cup \overline{C}$. \square

Finally, we show that given a constrained decision-function (F, C) , simply considering reasons of the total Boolean function F (and ignoring the constraint C), may actually supply strictly less succinct reasons.

Example 2. Continuing Example 1, note that the only sufficient reason for $F(0, 0) = 1$ is subsumed by a sufficient reason of $f_C(0, 0) = 1$, i.e., $(\neg X_1 \wedge \neg X_2)$ is subsumed by $\neg X_2$. Similarly, the only sufficient reason for $F(1, 1) = 1$ is subsumed by a sufficient reason of $f_C(1, 1) = 1$, i.e., $(X_1 \wedge X_2)$ is subsumed by X_1 . This accords with the intuition that, in light of the constraint $(X_1 \rightarrow X_2)$, reason X_1 is preferred to reason $(X_1 \wedge X_2)$. This is no accident: every reason using F is no more succinct than some reason

²Note that \mathbf{x} is a positive instance of f_C if and only if it is a positive instance of F and is in C . Thus, for a positive instance of F that satisfies the constraints, we can compare reasons using F with reasons using f_C .

using f_C , see Theorem 1. As this example shows, they can in fact be less succinct.

If t properly subsumes s then t is smaller than s (i.e., contains less literals), and it is not hard to find examples where every sufficient reason of $F(\mathbf{x}) = 1$ is much larger than every sufficient reason of $f_C(\mathbf{x}) = 1$. To do this, consider the constraint $X_1 \rightarrow (X_2 \wedge X_3 \cdots \wedge X_n)$: then every reason $X_1 \wedge X_2 \wedge \cdots \wedge X_n$ is subsumed by the reason X_1 .

Constraint-equivalent reasons. If one is interested in the meaning of a reason, and not its syntactic structure, then one should consider sufficient reasons up to logical-equivalence modulo the constraints. That is, terms t, s are *C-equivalent* (or simply, *constraint equivalent* when the constraint is understood), if $C \cap [s] = C \cap [t]$. For instance, if C is represented by $(X_1 \vee X_2) \wedge \neg(X_1 \wedge X_2)$ then $t = \neg X_1$ is *C-equivalent* to $s = X_2$, and thus s and t may be identified as the same reason in the presence of C .

Variations. Subtle changes in the definition of sufficient reasons result in radically different types of reasons. First, we have seen in the Examples that ignoring the constraints does not provide the most parsimonious reasons. Second, consider the variation in which, instead of asking for reasons of the partial function f_C , one uses the total function that maps \mathbf{x} to 1 if $F(\mathbf{x}) = 1$ and $\mathbf{x} \in C$, and to 0 otherwise. Although seemingly natural (indeed, why not assign a negative value to instances that do not satisfy the constraint), it is not hard to see, using Lemma 1, that this results in the *least parsimonious* reasons for F . On the other hand, using f_C , which amounts to assigning a positive value to instances that do not satisfy the constraint (as we will show in Proposition 1), results in the *most parsimonious* reasons. We find it striking that this change of perspective has such drastic changes on the parsimony of the produced reasons.

5 Computing sufficient reasons

In this section we discuss how to computationally find sufficient reasons in the presence of constraints. In particular, we show to reduce this to the unconstrained case.

Definition 6 (Computational problem). *Given a constrained decision-function (F, C) , and a positive instance \mathbf{x} of f_C , find all sufficient reasons for the decision $f_C(\mathbf{x}) = 1$.*

As usual (see the Preliminaries), we can think of the total function F and the set of instances C as Boolean formulas, say $F^1 = [\varphi]$ and $C = [\kappa]$ (we are agnostic about exactly how to represent these formulas until we discuss complexity and the experiments). The following proposition says that we can reduce the computational problem of the constrained case to the unconstrained case using the formula $(\kappa \rightarrow \varphi)$.

Proposition 1. *Suppose φ represents F and κ represents C . For a positive instance \mathbf{x} of f_C , the sufficient reasons of the decision $f_C(\mathbf{x}) = 1$ are exactly the sufficient reasons of the decision $G(\mathbf{x}) = 1$ where G is the total function represented by the Boolean formula $(\kappa \rightarrow \varphi)$.*

Proof. First, note that \mathbf{x} is a positive instance of G . Indeed, since $f_C(\mathbf{x}) = 1$ we know that $\mathbf{x} \in C \cap f^1$, i.e., $\mathbf{x} \models \kappa \wedge \varphi$, and thus also $\mathbf{x} \models \kappa \rightarrow \varphi$. Thus, it is sufficient to show that

a term t is an implicant of f_C iff it is an implicant of G . By definition, t is an implicant of f_C iff $[t] \subseteq (f_C)^1 \cup (f_C)^*$. But $(f_C)^1 = f^1 \cap C$ and $(f_C)^* = \overline{C}$ (Definition 4). On the other hand, t is an implicant of the total function G iff $[t] \subseteq G^1$. But $G^1 = \overline{C} \cup f^1$ by the law of distributivity. Thus $G^1 = (f_C)^1 \cup (f_C)^*$. \square

The significance of Proposition 1 is that it shows how to reuse algorithms and tools that are already developed for reasoning about total Boolean functions. Indeed, as long as the formulas κ, φ are represented in a language that allows one to form the conditional $\kappa \rightarrow \varphi$ formula in polynomial time in the sizes of κ, φ , we have a polynomial time reduction of the problem of finding reasons with constraints to those without. On the other hand, reasoning without constraints is a special case of reasoning with constraints, i.e., there is a trivial reduction in the other direction too, simply take $\kappa = \text{true}$. We summarise this important computational fact as follows:

Theorem 2. *Assume that formulas are represented in a formalism that allows one to form the conditional of two formulas in polynomial time. Then, the problem of finding all sufficient reasons for a decision with constraints is polynomial time interreducible with the problem of finding all sufficient reasons for a decision without constraints.*

Thus, if one uses representations that also allow one to compute sufficient reasons of *total* Boolean functions in polynomial time, then, by first applying the reduction in Theorem 2 one can find sufficient reasons for constrained decision-functions in polynomial time too.

We mention the two main approaches comprising the state of the art for computing principled reasons for total Boolean functions. First, (Shih, Choi, and Darwiche 2018) represent formulas using OBDDs, which do indeed support polynomial negation and conjunction. Their approach provides a polynomial time procedure for finding all sufficient reasons, using the fact that OBDDs support polynomial-time validity and entailment checking. To reuse their algorithm in our setting, simply run it on the OBDD representation of the formula $(\kappa \rightarrow \varphi)$. Second, (Ignatiev, Narodytska, and Marques-Silva 2019) take an agnostic view on the representation of formulas, and only require that the chosen representation allows polynomial time entailment checking. To reuse their approach in the presence of constraints, again, simply use it on formulas of the form $(\kappa \rightarrow \varphi)$ instead of φ .

Note that if a representation also allows (a) polynomial time validity checking, and (b) forming the conjunction of a term and formula in polynomial time, then one can decide if two terms are constraint-equivalent in polynomial time. Thus, if one is interested in computing reasons up to constraint-equivalence one can compute a set of representatives by, for instance, checking each pair of reasons for constraint-equivalence.

5.1 Illustration

To clarify the introduced concepts, we illustrate sufficient reasons on a complete synthetic example of a learnt classifier, inspired by an example in (Kisa et al. 2014).

Consider a tech-company that is shortlisting recent CS graduates for a job interview. The company considers candidates who took courses on Probability (P), Logic (L), Artificial Intelligence (A) or Knowledge Representation (K) during their studies. Suppose that the company uses data on candidates who were hired in the past to learn a threshold-based classifier, and let F be the associated *total* decision-function over $\mathbf{X} = \{L, K, P, A\}$ with $F^1 = \{(0011), (0110), (0111), (1100), (1101), (1110), (1111)\}$.³

Consider an instance $\mathbf{x} = (0011)$ corresponding to candidates that did not take L or K, but did take P and A. Note that $F(\mathbf{x}) = 1$, i.e., the classifier decides to grant such candidates interviews. What is the reason behind this decision? Table 2 gives the reasons which were computed using (Shih, Choi, and Darwiche 2018). We see that the only reason behind the decision of F for $\mathbf{x} = (0011)$ is $(\neg L \wedge P \wedge A)$, i.e., that the candidate did not take L, but did take P and A.

L	K	P	A	Reasons using F	using f_C
0	0	1	1	$(\neg L \wedge P \wedge A)$	$(\neg L \wedge A)$
0	1	1	1	$(\neg L \wedge P \wedge A), (K \wedge P)$	$(\neg L \wedge A), K$
1	1	0	0	$(L \wedge K)$	K
1	1	1	0	$(L \wedge K), (K \wedge P)$	K
1	1	1	1	$(L \wedge K), (K \wedge P)$	K

Table 2: All the positive instances \mathbf{x} that satisfy the constraints; Reasons for decisions $F(\mathbf{x}) = 1$; Reasons for decisions $f_C(\mathbf{x}) = 1$

Suppose, that a student’s enrolments must satisfy the following constraints C : a student must take P or L, ($P \vee L$); the prerequisite for A is P, ($A \rightarrow P$); the prerequisite for K is A or L, ($K \rightarrow (A \vee L)$). Reasons of the constrained decision-function f_C are given in Table 2. Note $(\neg L \wedge A)$ and K are not constraint-equivalent.

Consider the reason behind the decision $f_C(\mathbf{x}) = 1$ where $\mathbf{x} = (0011)$, i.e., $\neg L \wedge A$. This reason strictly subsumes the reason $\neg L \wedge P \wedge A$ used by the original (unconstrained) classifier F . This phenomenon, that for every positive instance \mathbf{x} in C , every sufficient reason of $F(\mathbf{x}) = 1$ is subsumed by some sufficient reason of $f_C(\mathbf{x}) = 1$, can be seen in all other rows of Table 2. This illustrates that our notion of sufficient reason (Definition 5) systematically eliminates such redundancies, a fact we formalised in Theorem 1.

6 Case Studies and Validation

In this section we validate our theory on constrained decision-functions learnt from binary data.⁴ We provide a prototype using a type of classifier that is often considered interpretable, i.e., decision trees. The purpose of the prototype is to provide a proof of concept that shows that by using constrained decision-functions: (1) we get no less succinct,

³The details of how such an F is learnt is not the focus of this paper, see, e.g., (Shih, Choi, and Darwiche 2018) for Bayesian networks, (Choi et al. 2019) for Neural Networks, and (Audemard, Koriche, and Marquis 2020) for Random Forests.

⁴Continuous data can be discretised, and discrete/categorical data can be binarised (Breiman et al. 1984).

and sometimes more succinct, reasons compared with the unconstrained setting; (2) we can seamlessly integrate two major types of constraints that can arise in AI, namely constraints due to pre-processing of data (e.g. one-hot, or other categorical, encodings), and semantic constraints that are inherent to the input domain.

Representation As discussed in Section 5 we can compute reasons by reducing to the unconstrained case. We reuse the algorithms in (Shih, Choi, and Darwiche 2018) by simply building an OBDD representing $\kappa \rightarrow \varphi$ (using the OBDD operations for complementation and disjunction), and pass this OBDD as input to their tool that computes sufficient reasons for a given instance.

Case Study 1. We used the dataset of Corticosteroid Randomization after Significant Head Injury (CRASH) trial (Collaborators et al. 2008) to predict the condition of a patient after a traumatic head injury. There are eleven clinically relevant input variables including demographics, injury characteristics and image findings, see (Zador, Sperrin, and King 2016). Six variables are categorical, and the rest are Boolean.⁵ Outcome variable *gos* 1 (0) indicates moderate or full recovery at 6 months versus death or severe disability.

Categorical variables are encoded using a one-hot encoding, which induces the constraint C as follows. For a categorical variable X , let D denote the set of Boolean variables corresponding to the set of categories of X . The corresponding constraint says that exactly one of the variables in D must be true. For example, variable *Eye* (shortened to *E*) has 4 categories, which we encode by the Boolean variables in $D_E = \{E_1, E_2, E_3, E_4\}$. The corresponding constraint is $\bigvee_i E_i \wedge \bigwedge_{i \neq j} \neg(E_i \wedge E_j)$, where i, j vary over $\{1, 2, 3, 4\}$. The constraint C is the conjunction of all such constraints, one for each categorical variable.

Following (Steyerberg et al. 2008) we base our example on 6945 cases with no missing values. RPART (seed: 25, train: 0.75, cp: 0.005) correctly classifies 75.69% of instances in the test set (ROC 0.77). Figure 2 shows the model.




Figure 2: RPART decision tree for Case Study 1.

Consider the instance \mathbf{x} that maps $A_1, E_1, M_5, V_2, P_1, OB, MD$ to 1, and the remaining four variables to 0. The decision-rule in the decision tree that explains why \mathbf{x} is positive is $E_1 \wedge P_1 \wedge \neg A_7 \wedge M_5 \wedge \neg A_6$ (size:5). There is one sufficient reason obtained from the decision-function F :

⁵Categorical are: Age(1-7), Eye(1-4) Motor(1-6), Verbal(1-5), Pupils(1-3). Boolean are: EC, PH, OB, SA, MD, HM.

$\neg A_6 \wedge \neg A_7 \wedge M_5 \wedge P_1$ (size:4). There are two sufficient reasons from f_C up to constraint-equivalence: (i) $A_1 \wedge M_5 \wedge P_1$ (size:3), (ii) $\neg A_6 \wedge \neg A_7 \wedge M_5 \wedge P_1$ (size:4).

Discussion of Case-Study 1. The explanation using the decision tree is strictly subsumed by the sufficient reason using F . This shows that decision-rules may not be the best reasons. Further, incorporating constraints resulted in having a smaller reason which would be completely missed if we just used F . Note that, as guaranteed by Theorem 1, the reason using F is subsumed by some reason using f_C , in fact it appears as reason (ii).

Note that reasons (i) and (ii) are not constraint-equivalent (and thus should be considered different reasons). Which reason should one prefer? On the one hand, (i) is more succinct, but on the other hand (ii) strictly constraint-subsumes (i), i.e., it applies to more instances. Without additional preferences there is no basis to prefer one over the other, and thus we report both of them

We remark that if one incorporated constraints by instead using the function represented by the formula $(\kappa \wedge \varphi)$, one would get one sufficient reason for this decision that is highly redundant in light of the constraint, i.e., $(A_1 \wedge E_1 \wedge M_5 \wedge V_2 \wedge P_1 \wedge \bigwedge_X \neg X)$ where the conjunction is over all the remaining variables $A_2, A_3, \dots, E_2, E_3, \dots$.

Finally, the histogram in Figure 3 compares the sizes of shortest reasons for F and f_C (omitting size 2 reasons which would dominate the graph). Note that the percentage of reasons using F increases with size, while those using f_C decreases with size.




Figure 3: Distribution of shortest reasons, restricted to instances without length ≤ 2 reasons (i.e., 5440 of 109120 instances). Percentages are rounded to the nearest decimal.

In summary, this case study validates empirically that with the constrained decision-functions, we may get smaller reasons compared to the unconstrained setting.

Case Study 2. To study semantic constraints, we used the Tic-Tac-Toe (TTT) Endgame dataset from the UCI machine learning repository (Dua and Graff 2017). This dataset contains the complete set of board configurations that result from X going first, until the game ends. The target concept is “player X has three-in-a-row”.

We binarise the dataset as in (Verwer and Zhang 2019). For each of the 9 board positions (labelled as in Table 3i.) introduce variables $V_{i,O}$ (resp. $V_{i,X}$) capturing whether or not O (resp. X) was placed in position i . We trained a classifier on this dataset using RPART (seed 1, train: 0.7, cp 0.01);

$$\text{i. } \begin{array}{|c|c|c|} \hline 0 & 1 & 2 \\ \hline 3 & 4 & 5 \\ \hline 6 & 7 & 8 \\ \hline \end{array} \quad \text{ii. } \begin{array}{|c|c|c|} \hline X & X & X \\ \hline O & & O \\ \hline \end{array} \quad \text{iii. } \begin{array}{|c|c|c|} \hline 01 & 01 & 01 \\ \hline 00 & 00 & 00 \\ \hline 10 & 00 & 10 \\ \hline \end{array}$$

Table 3: i. TTT board; ii. Positive instance; iii. Encoded instance (cell i is labelled by the values of $V_{i,O}V_{i,X}$).

with 93% accuracy for the test set (ROC 0.97), see Figure 4.




Figure 4: RPART decision tree for case study 2. We drop V and write, e.g., $4o$ instead of $V_{4,O}$ for readability.

Let F be the corresponding decision-function. In what follows we focus on sufficient reasons for the instance in Table 3iii. The sufficient reasons using F are given in Table 4.

0 - - 0 -	0 - - - -	- - - - 0 -	- - 0 - - -
1. - - 0 - -	2. - - 0 - 0 -	3. 0 - 0 - -	4. 0 - 0 - 0 -
- - 0 - - -	- - 0 - - -	- - 0 - - -	- - 0 - - -
01 - 1 01	01 - 1 - 1	-1 - 1 01	-1 01 - 1
5. - - - - -	6. - - - - 0 -	7. 0 - - - -	8. 0 - - - 0 -
- - 0 - - -	- - 0 - - -	- - 0 - - -	- - 0 - - -

Table 4: Sufficient reasons of using the decision-function F of the trained classifier, for the positive instance in Table 3iii.

Simple constraints for TTT The binarisation induces a constraint C that expresses that no position contains both an O and an X, although, unlike the one-hot-constraints in Case Study 1, it may have neither. So C is given by $\bigwedge_{0 \leq i \leq 8} \neg(V_{i,O} \wedge V_{i,X})$. Again, consider the positive instance in Table 3iii. The reasons for the decision using f_C include Reasons 1-4 in Table 4, as well as the following reason which strictly subsumes Reasons 5-8:

$$\begin{array}{ccc} -1 & -1 & -1 \\ \hline & & \\ - & - & - \\ \hline & 0 & - \end{array}$$

This shows that some reasons of F are redundant in light of the constraint C as witnessed, e.g., by the inclusion of the literals $\neg V_{0,O}$ and $V_{0,X}$ in reason 5.

More complex constraints: adding game rules Consider the constraint C' that includes C and adds the constraint that

says that the board is the result of valid play, i.e., that X moves first and players alternate moves. The additional constraint is $\bigvee_{S,T} (\psi_S \wedge \varphi_T)$ where S, T vary over all subsets of $U = \{0, 1, 2, \dots, 8\}$ such that $S \cap T = \emptyset$, and $0 \leq |S| - |T| \leq 1$, and ψ_S is $(\bigwedge_{i \in S} V_{i,X}) \wedge (\bigwedge_{i \in U \setminus S} \neg V_{i,X})$ and φ_T is $(\bigwedge_{i \in T} V_{i,O}) \wedge (\bigwedge_{i \in U \setminus T} \neg V_{i,O})$. The formula expresses that the set S of positions where X has played is disjoint from the set T where O has played, and that either there are the same number of moves, or X has played one more. Using $f_{C'}$, the sufficient reasons for the instance above include the following reason:

```
-- -- --
-- 00 00
-- 00 1-
```

This reason can be interpreted as follows: *in light of the constraint C'*, which says that the board is the result of a valid play, if positions 4,5,7 are blank and position 8 has an O, then player X must have won. This is indeed correct: player O could not have won since with 5 moves in the game player O can only move twice, and there could not be a draw because not all positions were filled yet.

Discussion of Case-Study 2. This case study illustrates how our framework can seamlessly take complex semantic constraints into account when producing reasons. Although one may be able to provide ad-hoc algorithms for incorporating specific constraints, our framework systematically and elegantly handles complex semantic constraints, such as the combinatorial constraint that captures valid play.

7 Related Work and Discussion

As discussed in the body, our theory generalises (Shih, Choi, and Darwiche 2018) and refines (Ignatiev, Narodytska, and Marques-Silva 2019) by incorporating domain constraints. We do not deal with constraints on the possible outputs of a classifier (Xu et al. 2018). To provide sufficient reasons in the presence of domain constraints, we show how to reduce the constrained case to the unconstrained case, thus allowing one to reuse existing symbolic algorithms and tools (Shih, Choi, and Darwiche 2018). Other approaches for handling the unconstrained case include heuristics (Ribeiro, Singh, and Guestrin 2016; Lakkaraju et al. 2019; Iyer et al. 2018) and a combination of heuristic and abductive reasoning (Ignatiev, Narodytska, and Marques-Silva 2019).

The ML literature has techniques for producing (post-hoc local) rule-based explanations which are similar in spirit to the logic-based method of this paper. Notably, the *anchors* of (Ribeiro, Singh, and Guestrin 2018) are analogous to implicants. That work is probabilistic, e.g., it works directly on a probabilistic model while our method works on Boolean functions representing a possibly probabilistic model. That work aims to optimise the coverage of anchors (i.e., the probability that the anchor applies to a random instance), which is not an analogue of prime-implicant, and thus potentially misses out valid explanations (indeed, that work has no analogue of subsumption).

None of the works cited above explicitly handle constraints, which is the main focus of our work.

8 Discussion

The crux of this paper shows how to handle constraints in a principled manner, and establishes that ignoring constraints could result in unnecessarily long/complex reasons, as well as missing some reasons altogether.

A general critique of the prime-implicant based approach to reasoning is that reasons may become too large to comprehend when the number of variables is large. Notice that our method is a step towards improving this problem in the presence of constraints. If the shortest reason in the presence of constraints is still too large to comprehend, not taking constraints into account may result in reasons that are even larger and even harder to comprehend.

To validate the claim that using constraints results in no longer, and sometimes shorter reasons, in the two case studies, we compared (with equivalence, subsumption and constraint-equivalence tests), every reason obtained from decision-function F with that of f_C and observed that while adding constraints may decrease or increase the *number* of reasons, it never increases the *size* of the shortest reasons (a fact that is guaranteed by Theorem 1).

In cases of multiple (constraint-inequivalent) reasons for a decision (even amongst the shortest ones), we do not supply a way to pick one reason over another, a challenging problem (Lakkaraju et al. 2019). Indeed, preferring one reason over another would require *additional assumptions* about preferred reasons, e.g., favouring succinctness over generality (Miller 2019).

Our approach to handling constraints is general enough that it can be applied, in principle, to any type of classifier, no matter how it was learnt, and no matter how it is represented. As a proof-of-concept, we illustrated this theoretical advance in the case studies using decision trees. Since small decision trees are often considered interpretable, they also allow one to “read-off” reasons from their branches in order to compare with our most-parsimonious reasons. In fact, using decision trees yielded the observation in Case Study 1 that the decision-tree branches may not supply the most succinct reasons, a point worthy of further study.

We remark that since our approach reduces the constrained case to the unconstrained case to compute reasons, any advance in the efficiency of tools for solving the latter will yield benefits for the former.

Our work opens up applications that are currently only available in the unconstrained setting, including the study of classifier bias and counterfactual decisions (Darwiche and Hirth 2020).

References

- Audemard, G.; Koriche, F.; and Marquis, P. 2020. On tractable XAI queries based on compiled representations. In *KR*.
- Breiman, L.; Friedman, J.; Stone, C. J.; and Olshen, R. A. 1984. *Classification and regression trees*. CRC press.
- Choi, A.; Shi, W.; Shih, A.; and Darwiche, A. 2019. Compiling neural networks into tractable Boolean circuits. *AAAI VNN*.

- Collaborators, M. C. T.; et al. 2008. Predicting outcome after traumatic brain injury: practical prognostic models based on large cohort of international patients. *BMJ*, 336(7641).
- Coudert, O. 1994. Two-level logic minimization: an overview. *Integration*, 17(2).
- Darwiche, A. 2020. Three Modern Roles for Logic in AI. In *PODS*.
- Darwiche, A.; and Hirth, A. 2020. On The Reasons Behind Decisions. In *ECAI*.
- Darwiche, A.; and Marquis, P. 2002. A knowledge compilation map. *Journal of Artif. Intell. Research*, 17: 229–264.
- Dua, D.; and Graff, C. 2017. UCI Machine Learning Repository.
- Ignatiev, A.; Narodytska, N.; and Marques-Silva, J. 2019. Abduction-based explanations for machine learning models. In *AAAI*.
- Iyer, R.; Li, Y.; Li, H.; Lewis, M.; Sundar, R.; and Sycara, K. 2018. Transparency and explanation in deep reinforcement learning neural networks. In *AAAI*.
- Kisa, D.; Van den Broeck, G.; Choi, A.; and Darwiche, A. 2014. Probabilistic sentential decision diagrams: Learning with massive logical constraints. In *ICML LTPM*.
- Lakkaraju, H.; Kamar, E.; Caruana, R.; and Leskovec, J. 2019. Faithful and customizable explanations of black box models. In *AAAI*.
- McCluskey, E. J. 1956. Minimization of Boolean functions. *Bell Syst.*, 35(6): 1417–1444.
- Miller, T. 2019. Explanation in artificial intelligence: Insights from the social sciences. *Artif. Intell.*, 267.
- Quine, W. V. 1952. The problem of simplifying truth functions. *Am Math Mon*, 59(8).
- Ribeiro, M. T.; Singh, S.; and Guestrin, C. 2016. “Why should I trust you?” Explaining the predictions of any classifier. In *SIGKDD*.
- Ribeiro, M. T.; Singh, S.; and Guestrin, C. 2018. Anchors: High-Precision Model-Agnostic Explanations. *AAAI*.
- Shih, A.; Choi, A.; and Darwiche, A. 2018. A symbolic approach to explaining bayesian network classifiers. In *IJCAI*.
- Steyerberg, E. W.; Mushkudiani, N.; Perel, P.; Butcher, I.; Lu, J.; McHugh, G. S.; Murray, G. D.; Marmarou, A.; Roberts, I.; Habbema, J. D. F.; et al. 2008. Predicting outcome after traumatic brain injury: development and international validation of prognostic scores based on admission characteristics. *PLoS medicine*, 5(8).
- Verwer, S.; and Zhang, Y. 2019. Learning optimal classification trees using a binary linear program formulation. In *AAAI*.
- Xu, J.; Zhang, Z.; Friedman, T.; Liang, Y.; and Broeck, G. 2018. A semantic loss function for deep learning with symbolic knowledge. In *ICML*.
- Zador, Z.; Sperrin, M.; and King, A. T. 2016. Predictors of outcome in traumatic brain injury: new insight using receiver operating curve indices and Bayesian network analysis. *PLoS one*, 11(7).