

Causal Intervention for Subject-deconfounded Facial Action Unit Recognition

Yingjie Chen¹, Diqi Chen², Tao Wang^{1*}, Yizhou Wang¹, Yun Liang¹

¹ School of Computer Science, Peking University, Beijing China

² Advanced Institute of Information Technology (AIIT), Peking University, Hangzhou, China

chenyingjie@pku.edu.cn, dqchen@aiit.org.cn, wangtao@pku.edu.cn, yizhou.wang@pku.edu.cn, ericlyun@pku.edu.cn

Abstract

Subject-invariant facial action unit (AU) recognition remains challenging for the reason that the data distribution varies among subjects. In this paper, we propose a causal inference framework for subject-invariant facial action unit recognition. To illustrate the causal effect existing in AU recognition task, we formulate the causalities among facial images, subjects, latent AU semantic relations, and estimated AU occurrence probabilities via a structural causal model. By constructing such a causal diagram, we clarify the causal effect among variables and propose a plug-in causal intervention module, CIS, to deconfound the confounder *Subject* in the causal diagram. Extensive experiments conducted on two commonly used AU benchmark datasets, BP4D and DISFA, show the effectiveness of our CIS, and the model with CIS inserted, CISNet, has achieved state-of-the-art performance.

Introduction

With the proliferation of facial behavior analysis in real-world application scenarios such as online education and driver safety assistance, facial action unit recognition has attracted increasing research interest as a fundamental task in the field of affective computing. According to Facial Action Coding System (FACS) (Friesen and Ekman 1978), facial action units (AUs), defined as the combinations of facial muscle movements, can describe almost all facial behaviors, which is essential for fine-grained facial behavior analysis. In recent years, deep learning has proved its efficacy and efficiency in facial action unit recognition task (Cui et al. 2020; Chen et al. 2021b; Yang et al. 2021; Song et al. 2021b), but there is still room for improvement since some inherent nature of AU has not been fully exploited.

AUs are not independent of each other. On the one hand, AUs usually do not occur alone when humans express certain emotions, and thus some combinations of AUs, which pertain to displayed emotions, can be frequently observed, e.g. AU6 (Cheek Raiser) and AU12 (Lip Corner Puller) tend to appear together and form facial expression *happiness* (Ekman 1992). On the other hand, there are strict co-occurrence and mutual exclusion among AUs due to the structural con-




Figure 1: Illustration of subject variation problem. The AU semantic relations embedded in the facial images of the four subjects vary due to the differences among their customs of expressing happiness. Using samples collected from the first three subjects for training may cause the learned latent AU semantic relations containing subject-specific ones (yellow lines) in addition to universal ones (blue lines). When encountering a new subject (*Subject 4*) in the inference stage, differences between *Subject 4*'s specific AU relations and the learned ones will lead to prediction bias.

straints brought by facial anatomy, e.g. AU22 (Lip Funneler) and AU23 (Lip Tightener) cannot appear simultaneously since they are all related to facial muscle *Orbicularis oris*, and it is difficult to make AU9 (Nose Wrinkler) without the presence of AU6 (Cheek Raiser) due to muscular synergy (Zarins 2018). Although the presence or absence of each AU can be mainly inferred from facial appearance changes, it can also be partially inferred based on the states of other AUs. Therefore, an accurate AU recognition model

*Corresponding author.

Copyright © 2022, Association for the Advancement of Artificial Intelligence (www.aaai.org). All rights reserved.

captures not only low-level facial appearance features but also high-level semantic relations among AUs.

Recent works have made progress in capturing high-level AU semantic relations in an implicit way (Corneanu, Madadi, and Escalera 2018; Niu et al. 2019) by exploiting correlations between AUs via probabilistic graphic models or in an explicit way (Li et al. 2019; Shao et al. 2020) by constructing an AU semantic graph according to statistics of the training data, and both kinds of works have achieved more accurate AU recognition. Although these methods can make use of priors contained in the training data, they all suffer from prediction bias while being applied to samples of new subjects. This is known as subject variation problem, which makes it challenging for AU recognition model to generalize across subjects. Although previous works have noticed that subject variation problem exists in facial action unit recognition task, as far as we know, there have been few works focusing on answering the whys and wherefores.

We argue that the prediction bias caused by subject variation problem mainly due to the fact that the latent AU semantic relations vary among subjects. As shown in Fig. 1, subjects in all facial images are expressing the facial expression *happiness*, which is composed of AU6 and AU12 as mentioned in (Ekman 1992). Thus the co-occurrence of AU6 and AU12 is one universal AU semantic relation shared by all subjects. However, AU semantic relations embedded in each facial image vary among subjects, which means that in addition to the universal relation, there are subject-specific AU semantic relations due to the differences in subjects' customs of expressing emotions. *E.g.* *Subject 1* tends to raise her inner brows (AU1) while smiling and *Subject 2* tends to raise his whole brows (AU1 and AU2) and laugh with the drop of his jaw (AU26). If we train a model using samples of the first three subjects, the model will learn a set of latent AU semantic relations containing both the universal ones and the subject-specific ones. When applying the model to *Subject 4*, the subject-specific relations of the training subjects may lead to prediction bias on AU1, AU2, or AU26.

So far, we can see that the prediction bias of AU recognition is mainly caused by the differences among subjects' customs of expressing emotions. *Subject* can be essentially regarded as a confounder, which misleads AU recognition model to learn subject-specific AU semantic relations from subjects in the training data and thus causes prediction bias while applying the model to a new subject for inference. For clarity, we denote the input facial images as X and the predicted AU occurrence probabilities as Y , and an AU recognition model aims to approximate $P(Y|X)$ as much as possible. However, as mentioned above, $P(Y|X)$ may lead to a biased AU recognition model since it may learn subject-specific AU relations that are not shared by new subjects. For example, $P(Y|X)$ would learn the relation that AU2 and AU12 tend to co-occur when using samples of *Subject 1* in Fig. 1 for training, which is actually a subject-specific relation of *Subject 1* which is not suitable for others, and $P(Y|X)$ would mistakenly reduce the importance of the relation that AU6 and AU12 co-occur and form the facial expression *happiness* when using facial images of *Subject 3* for training, since *Subject 3* may have physical difficulty in

contracting facial muscles related to the occurrence of AU6. To relieve subject variation problem, we propose a method to learn the universal AU semantic relations by making our model approximate $P(Y|do(X))$ instead of $P(Y|X)$, where *do*-operation denotes the pursuit of the causality between the cause X and the effect Y without the confounding effect caused by the confounder *Subject*.

To this end, we formulate subject variation problem by constructing a causal diagram to analyze the causalities among facial images, subjects, latent AU semantic relations, and estimated AU occurrence probabilities. Our causal inference framework not only fundamentally explains how subject-specific AU semantic relations hurt the performance of AU recognition models, but also provides a solution by removing the effect caused by confounder *Subject*. Based on the causal model, a plug-in causal intervention module called **CIS** is proposed to deconfound *Subject* via back-door adjustment (Pearl, Glymour, and Jewell 2016).

Our main contributions are listed as:

- We formulate subject variant problem in AU recognition using an AU causal diagram to explain the whys and wherefores. To the best of our knowledge, this is the first work to explain this problem with the help of causal inference theory and make attempt to remove the effect caused by subject variation via causal intervention.
- Based on our causal diagram, we proposed a plug-in causal intervention module, CIS, which could be inserted into advanced AU recognition models for removing the effect caused by confounder *Subject*.
- Extensive experiments on two widely used AU benchmark datasets, BP4D and DISFA, demonstrate that the proposed CIS can boost various AU recognition methods to new state-of-the-art.

Related Work

Facial Action Recognition

In recent years, research on facial action unit recognition has seen great achievements. The rise of deep learning has raised the performance of AU recognition to a new level. Considering the locality of AUs, methods such as (Zhao, Chu, and Zhang 2016; Li, Abtahi, and Zhu 2017; Li et al. 2017; Song et al. 2021a; Chen et al. 2021a) make attempt to learn better facial appearance features by emphasizing important local facial regions. Zhao et al. (Zhao, Chu, and Zhang 2016) proposed Deep Region and Multi-label Learning (DRML), which employs a region layer to induce important facial regions and force the learned weights to capture structural information of the face. Considering the semantic relations among AUs, some works (Wang et al. 2013; Walecki et al. 2017) make efforts in modeling such relations via probabilistic graphical models or graph neural networks. Wang et al. (Wang et al. 2013) introduced a restricted Boltzmann machine to model the action units, thereby capturing not only local but also global AU dependencies. Li et al. (Li et al. 2019) investigated how to integrate the semantic relationship propagation between AUs to enhance the feature representation of facial regions, and proposed an AU semantic relationship embedded representation learning (SRERL)




Figure 2: Illustration of our AU causal diagram.

framework. However, these works ignore subject variation problem and thus the obtained AU recognition models suffer from the subject-related prediction bias.

As for subject variation problem, works such as (Chen et al. 2013) provide a solution for enhancing the generalizability of AU recognition model by training personalized AU classifiers for each subject and works such as (Zen et al. 2016; Wang and Wang 2018) make attempt to relieve the subject-related prediction bias through domain adaptation. Although these works have realized that the data distribution of training subjects differs from that of unseen subjects, they are still based on the assumption that the data distribution of source and target domains shares some similarities. In contrast, we formulate the causalities among variables in AU recognition task via a structural causal model to answer the whys and wherefores of subject variation problem and provide a solution based on causal intervention.

Causal Inference in Computer Vision

Causal inference (Pearl et al. 2000; Rubin 2005) has been gradually applied to computer vision tasks in recent years, such as long-tailed classification (Tang, Huang, and Zhang 2020), weakly-supervised semantic segmentation (Zhang et al. 2020), few-shot learning (Yue et al. 2020), and class-incremental learning (Hu et al. 2021). Causal inference empowers models the capability to consider the causal effect that naturally exists in a task and disentangle direct effect and indirect effect. There are two ways for causal inference: one is Pearl’s structural causal model (Pearl et al. 2000), and the other is the potential outcome framework proposed by Robins and Greenland (Rubin 2005), in terms of which they set out to express their conception of confounding. The advantage of structural causal model is that it shows the causal and effect among several variables in the form of a causal diagram, which is more intuitive and conducive for analysis, and thus we choose to follow the first way in our work.

Methodology

AU Causal Diagram

To answer the whys and wherefores of subject variation problem, we use a structural causal model (Pearl et al. 2000) to illustrate the causalities among variables in AU recognition models. As shown in Fig. 2, there are four variables involved in our AU causal diagram, which are facial images X , subjects S , latent AU semantic relations R , and estimated AU occurrence probabilities Y , and the causalities among

them are formulated via causal links, *i.e.* the direct edges in the causal diagram, each of which denotes the causality between two nodes, *i.e.* cause \rightarrow effect. Causalities in our causal diagram are described in detail.

$S \rightarrow X$. Subjects’ customs of expressing emotions lead to subject-specific facial expressions in the facial images recorded while they are expressing emotions. In this way, S determines what facial expressions appear on a subject’s face, and for the same kind of facial expression, the facial appearance changes differ among subjects in a subtle but not negligible way.

$S \rightarrow R \leftarrow X$. R denotes latent AU semantic relations, which consist of both the universal AU semantic relations and the subject-specific AU semantic relations. The universal ones are determined by the facial anatomical basis of AUs, which is universal to all subjects. However, due to the custom differences in the way that subjects express emotions, not only the universal AU semantic relations but also the subject-specific AU semantic relations can be observed in the facial images recorded from one subject. The latent AU semantic relations embedded in facial images are reflected by the causal link from X to R . The learned AU semantic relations are embedded in a pre-trained AU recognition model, which contains subject-specific ones and can be reflected by the causal link from S to R .

$X \rightarrow Y \leftarrow R$. Conventional AU recognition models aim to estimate AU occurrence probabilities Y as precisely as possible. From our causal diagram, we can see that Y is the effect of two causal paths, which are $X \rightarrow Y$ and $R \rightarrow Y$. The first causal path denotes that an AU recognition model estimates Y based on the facial appearance features extracted from the input facial image, and the second one denotes that the learned latent AU semantic relations embedded in the model influence the estimated Y by making use of priors from the training data. In other words, conventional AU recognition models which approximate $P(Y|X)$ learn a set of latent AU semantic relations R from the training data, which can be regarded as a kind of priors influencing the estimation results Y in an implicit way. Although the effect brought by R may introduce priors from the training data for better estimation when the facial appearance features are not sufficient enough to determine the states of certain AUs, it is confounded by subjects for training and thus may mistakenly associate or disassociate certain AUs.

Causal Intervention via Back-door Adjustment

To remove the adverse effect brought by confounder S and obtain a model which estimates Y only based on what’s in X , *i.e.* facial appearance features of the input facial images, we propose to intervene X by applying *do*-operator to variable X . The *do*-operator erases all the arrows that come into X , and in this way it prevents any information about X from flowing in the noncausal direction. In this way, the causal link from S to X is cut-off, and we obtain an AU recognition model approximating $P(Y|do(X))$ instead of $P(Y|X)$.

The straightforward way to intervene X is conducting a randomized controlled trial by collecting any facial image




Figure 3: Overview. First, a facial image is fed into a backbone network for feature extraction. Instead of directly using the extracted feature f_{cur} for classification, we put f_{cur} into the proposed CIS module for causal intervention on *Subject*, i.e. approximation of $P(Y|do(X))$. In CIS module, the output of *Approx. R* r_{cur} and f_{cur} are further fed into a linear layer separately and concatenated as the input of a classifier for AU prediction. The key component in CIS module is the approximation of R , which involves three parts for calculation, confounder attentions, a confounder dictionary, and confounder priors.

of any subject, and in this circumstance, $P(Y|X)$ equals $P(Y|do(X))$. Since such kind of intervention is impossible due to the infinity of the number of subjects and facial images, we apply the back-door adjustment formula as described in (Pearl, Glymour, and Jewell 2016) to obtain $P(Y|do(X))$ in a simpler and feasible way. To do this, we first estimate the effect at each stratum of the deconfounder in the intervention and then compute a weighted average of those strata, where each stratum is weighted according to its proportion. To be specific, the deconfounder is *Subject* in our case, and we first estimate the causal effect for each subject in the training data and then estimate the average causal effect by computing a weighted average based on the proportion of each subject's facial images in the training data.

As described above, the effect of the intervention, $P(Y|do(X))$, can be formulated as follows:

$$P(Y|do(X)) = \sum_s P(Y|X, R = f(X, s))P(s), \quad (1)$$

where $f(\cdot)$ is a function with X and each s as independent variables and R as dependent variable. As S is no longer correlated with X , the causal intervention makes X have a fair opportunity to incorporate every subject s into the estimation of Y , based on the proportion of each s in the whole. After that, the causal link from S to X is cut off, which al-

lows the causal effect $X \rightarrow Y$ free from the effect of S .

CISNet Architecture

Overview As shown in Fig. 3, our CISNet takes one facial image X as input, and a backbone network is first applied to the facial image for feature extraction. Then the extracted feature $f_{cur} \in \mathbb{R}^{d_{in}}$ is passed through our plug-in causal intervention module, i.e. CIS module, for subject deconfounding. After that, the output of CIS module is fed into a classifier for AU recognition. By given a threshold τ , the estimated AU occurrence probabilities as the classifier's output are processed as the final binary AU prediction results.

CIS Module We propose a plug-in causal intervention module, named CIS, for the approximation of the theoretical back-door adjustment formula. As the calculation of Eq. 1 requires the forward steps for each pair of X and s , which cost a lot. Thanks to the Normalized Weighted Geometric Mean (NWGM) mentioned in (Xu et al. 2015), which allows us to approximate the above expectation in feature-level. In this way, our aim turns out to be computing Eq. 2.

$$P(Y|do(X)) \stackrel{\text{NWGM}}{\approx} P(Y|X, R = \sum_s f(X, s)P(s)). \quad (2)$$

We apply a linear model to approximate the conditional probability, i.e. the probability of Y under the conditions X

and R , as shown in Eq. 3:

$$P(Y|do(X)) = W_X f_{\text{cur}} + W_S r_{\text{cur}}, \quad (3)$$

where $W_X \in \mathbb{R}^{d_{\text{out}} \times d_{\text{in}}}$ and $W_S \in \mathbb{R}^{d_{\text{out}} \times d_{\text{in}}}$ are learnable weight parameters, $r_{\text{cur}} \in \mathbb{R}^{d_{\text{in}}}$ is an approximation of R .

The key component in CIS module is the calculation of r_{cur} , which involves three parts, confounder attentions, a confounder dictionary, and confounder priors. For confounder S , since we cannot collect samples from all the subjects, we approximate S as a fixed confounder dictionary $S = [s_1, s_2, \dots, s_N]$, where N is the total number of subjects contained in the training data and each $s_i \in \mathbb{R}^{d_{\text{in}}}$ is the subject prototype of *Subject i*. To compute subject prototypes, we maintain one feature memory bank for each subject to store latent features, *i.e.* the output of the backbone network $\{f_k^i \in \mathbb{R}^{d_{\text{in}}} \}_{k=1}^{N_i}$, for training samples of the certain subject, where N_i is the number of samples for *Subject i*. And in this way, s_i is computed as $\frac{1}{N_i} \sum_{k=1}^{N_i} f_k^i$. Our confounder dictionary is updated as the end of each epoch, and confounder prior $P(s_i)$ can be computed as the ratio of the number *Subject i*'s samples to the total number of samples for training. We approximate R as a weighted aggregation of all subject prototypes, as shown in Eq. 4:

$$r_{\text{cur}} = \sum_{i=1}^N \alpha_i s_i P(s_i), \quad (4)$$

where α_i is a confounder attention for s_i in the confounder dictionary with specific f_{cur} , which is computed using Scaled Dot-Product Attention (Vaswani et al. 2017) as shown in Eq. 5.

$$\alpha_i = \text{softmax} \left(\frac{(W_Q f_{\text{cur}})^T (W_K s_i)}{\sqrt{d_m}} \right), \quad (5)$$

where $W_Q \in \mathbb{R}^{d_m \times d_{\text{in}}}$ and $W_K \in \mathbb{R}^{d_m \times d_{\text{in}}}$ are learnable weight parameters.

Loss Function To relief data imbalance problem in AU recognition, we apply an adaptive loss function for training:

$$\mathcal{L} = - \sum_{i=1}^C [(1 - \mu_i) p_i \log \hat{p}_i + \mu_i (1 - p_i) \log (1 - \hat{p}_i)], \quad (6)$$

where C is the number of AUs, p_i is the ground-truth binary label for the i^{th} AU, and \hat{p}_i denotes the estimated occurrence probability of the i^{th} AU. The occurrence frequency of the i^{th} AU in the training set, denoted by μ_i , is maintained adaptively to weight the loss of each AU.

Experiments

Datasets and Metrics

Two benchmark datasets, BP4D (Zhang et al. 2014) and DISFA (Mavadati et al. 2013), are used in our experiments. BP4D involves 23 female and 18 male young adults. Each subject finishes 8 tasks, and 324 videos are captured, resulting in around 140,000 images. Frame-level AU annotations

are provided by two FACS coders independently. DISFA involves 26 adult subjects, who are asked to watch videos to obtain their spontaneous facial behaviors. AUs are coded manually by a FACS coder with intensity labels within a scale of 0 to 5 for each frame. Frames with AU intensity labels greater than 1 are selected as positive samples.

We conduct a subject-exclusive 3-fold cross-validation on each dataset, following the experiment settings mentioned in (Li et al. 2017, 2019; Song et al. 2021b) for a fair comparison. We evaluate the proposed method using F1-score, which is commonly used in determining the harmonic mean between precision and recall.

Implementation Details

For each input image, we use Dlib (King 2009) to detect facial landmarks and align the image according to the computed coordinates of eye centers, crop the facial region, and resize the cropped face to 256×256 . We employ ResNet34 (He et al. 2016) without the final linear layer ($512 \rightarrow 64, 64 \rightarrow C$) as our backbone network, and two linear layers act as the classifier. d_{in} and d_{out} are set to 512 and d_m is set to 256. Before training, we set N feature memory banks to save features for each subject in the training data, and N equals 27, 27, 28 for each fold respectively. In the first epoch, we stop the back-propagation step to initialize our confounder dictionary. Threshold $\tau = 0.5$ is used for the binarization of the predicted probabilities $\hat{p} \in \mathbb{R}^C$.

CISNet is implemented on PyTorch (Paszke et al. 2017) platform. We use Stochastic gradient descent (SGD) with momentum of 0.9 and weight decay of 0.0005 as the optimizer. Learning rate is set to 0.001 and batch size is set to 4. The number of training epochs is set to 15, and early stopping strategy is used for training. Models are trained on one NVIDIA Tesla V100 16GB GPU.

Comparison with State-of-the-art Methods

We compare the performance of CISNet with the previous state-of-the-art methods including DRML (Zhao, Chu, and Zhang 2016), EAC-Net (Li et al. 2017), ROI-Net (Li, Abtahi, and Zhu 2017), DSIN (Corneanu, Madadi, and Escalera 2018), JAA-Net (Shao et al. 2018), LP-Net (Niu et al. 2019), SRERL (Li et al. 2019), UGN-B (Song et al. 2021a) and HMP-PS (Song et al. 2021b).

As shown in Table 1, CISNet outperforms all the compared methods in terms of average F1-score on BP4D. Specifically, CISNet achieves an average F1-score of 64.3%, which outperforms HMP-PS with the second-highest F1-score by 0.9%. Table 2 shows the experimental results on DISFA. CISNet outperforms HMP-PS by a large margin with the highest average F1-score of 64.7%. It is worth mentioning that the proposed CIS module can be inserted into almost all frame-based AU recognition models suffering from subject variation problem for causal intervention.

Ablation Study

Our ablation studies aim to answer: **Q1.** How does CIS module perform in models with different backbone networks? **Q2.** Is the number of training subjects the larger the better?

Method	AU1	AU2	AU4	AU6	AU7	AU10	AU12	AU14	AU15	AU17	AU23	AU24	Avg.
DRML	36.4	41.8	43.0	55.0	67.0	66.3	65.8	54.1	33.2	48.0	31.7	30.0	47.7
EAC-Net	39.0	35.2	48.6	76.1	72.9	81.9	86.2	58.8	37.5	59.1	35.9	35.8	55.6
ROI-Net	36.2	31.6	43.4	77.1	73.7	85.0	87.0	62.6	45.7	58.0	38.3	37.4	56.3
DSIN	51.7	40.4	56.0	76.1	73.5	79.9	85.4	62.7	37.3	62.9	38.8	41.6	58.9
JAA-Net	47.2	44.0	54.9	77.5	74.6	84.0	86.9	61.9	43.6	60.3	42.7	41.9	60.0
LP-Net	43.4	38.0	54.2	77.1	76.7	83.8	87.2	63.3	45.3	60.5	48.1	54.2	61.0
SRERL	46.9	45.3	55.6	77.1	78.4	83.5	87.6	63.9	52.2	63.9	47.1	53.3	62.9
UGN-B	54.2	46.4	56.8	76.2	76.7	82.4	86.1	64.7	51.2	63.1	48.5	53.6	63.3
HMP-PS	53.1	46.1	56.0	76.5	76.9	82.1	86.4	64.8	51.5	63.0	49.9	54.5	63.4
CISNet	54.8	48.3	57.2	76.2	76.5	85.2	<u>87.2</u>	66.2	50.9	65.0	47.7	56.5	64.3

Table 1: F1-score (%) for 12 AUs reported by the proposed method and the state-of-the-art methods on BP4D dataset. The best and second results are indicated using bold and underline, respectively.

Method	AU1	AU2	AU4	AU6	AU9	AU12	AU25	AU26	Avg.
DRML	17.3	17.7	37.4	29.0	10.7	37.7	38.5	20.1	26.1
EAC-Net	41.5	26.4	66.4	50.7	8.5	89.3	88.9	15.6	48.5
DSIN	42.4	39.0	68.4	28.6	46.8	70.8	90.4	42.2	53.6
JAA-Net	43.7	46.2	56.0	41.4	44.7	69.6	88.3	58.4	56.0
LP-Net	29.9	24.7	<u>72.7</u>	46.8	<u>49.6</u>	72.9	<u>93.8</u>	<u>65.0</u>	56.9
SRERL	<u>45.7</u>	47.8	<u>59.6</u>	47.1	<u>45.6</u>	73.5	<u>84.3</u>	<u>43.6</u>	55.9
UGN-B	43.3	<u>48.1</u>	63.4	49.5	48.2	72.9	90.8	<u>59.0</u>	59.4
HMP-PS	38.0	<u>45.9</u>	65.2	<u>50.9</u>	50.8	76.0	93.3	67.6	<u>61.0</u>
CISNet	48.8	50.4	78.9	51.9	47.1	<u>80.1</u>	95.4	<u>65.0</u>	64.7

Table 2: F1-score (%) for 8 AUs reported by the proposed method and the state-of-the-art methods on DISFA dataset. The best and second results are indicated using bold and underline, respectively.

A1. Effectiveness of CIS on Different Backbones To better elucidate the effectiveness of CIS, we evaluate the improvement brought by CIS module inserted in models with different backbone networks including ResNet18, ResNet34 and ResNet50 (He et al. 2016) in Table 3. By inserting CIS into models with different backbone networks, we can observe significant and consistent improvements compared with the corresponding models w/o CIS, which owes to the reduction of prediction bias via deconfounding of *Subject*.

A2. Impact of the Number of Training Subjects Considering that back-door adjustment is based on the premise of sufficient data, we conduct experiments to study the impact of the number of training subjects by training a baseline model (w/o CIS) and CISNet (w/ CIS), which only differ in terms of CIS module. Fig. 4 shows that with the number of training subjects increasing, the performance of two models




Figure 4: Impact of the number of training subject.

on the same test data increases accordingly, but the increasing rate of CISNet is slightly higher than the baseline model, which illustrates the importance of a sufficient number of training subjects and the potential of performance improvement brought by CIS module with more training subjects.

Qualitative Results

To better explain the mechanism of CIS module, we provide qualitative results in this section to answer: **Q3**. Does CIS module truly assist the model to estimate Y only based on what's in X ? **Q4**. Does the AU representations extracted by our CISNet invariant to subjects? **Q5**. What's the differences between models approximating $P(Y|X)$ or $P(Y|do(X))$?

Table 3: Effectiveness of CIS on different backbones.

BP4D			DISFA			
w/o CIS	w/ CIS	Δ	w/o CIS	w/ CIS	Δ	
ResNet18	59.9	63.8	3.9↑	57.3	63.5	6.2↑
ResNet34	60.6	64.3	3.7↑	57.4	64.7	7.3↑
ResNet50	61.3	63.6	2.3↑	56.2	65.0	8.8↑




Figure 5: PCC among AUs for different subjects. From left to right, PCC matrices are computed based on the ground-truth AU labels, predicted ones using CISNet (w/ CIS), and predicted ones using the baseline model (w/o CIS), respectively. Numbers under PCC heatmaps are cosine similarities between themselves and the corresponding ground-truth.

A3. Visualization of PCC Heatmaps To illustrate that CIS module acts as a causal intervention on *Subject* and makes the model estimate Y only based on X without unnecessary or even harmful prior from the training data, we visualize the Pearson Correlation Coefficient (PCC) matrices computed based on ground-truth or predicted AU labels for each subject. Fig. 5 shows that the PCC heatmaps computed for the baseline model and CISNet differ from each other, and the PCC heatmaps for CISNet are more similar to the PCC heatmaps computed based on the ground-truth AU labels according to their cosine similarities, which demonstrates that by using CIS to deconfound *Subject*, the model is endowed with the capability to focus on what's in X from a new subject without the confounding of latent subject-specific AU semantic relations implied in the training data.

A4. Visualization of Representations To show that the representations learned by CISNet are more subject-invariant, we insert C spatial-attention layers (Zhao and Wu 2019) between the backbone network and the classifier to obtain AU-specific features, where C is the number of AUs, and use t-SNE (van der Maaten and Hinton 2008) for visualization. As shown in Fig. 6, in the visualization result of the baseline model, features with the same subject labels form small clusters in a cluster corresponding to a certain AU. As for CISNet, features with the same subject labels are more dispersed in a cluster corresponding to a certain AU, which illustrates that the representations learned by CISNet are more invariant to subjects.

A5. Case Study We visualize the estimated AU occurrence probabilities of several samples to show the differences between $P(Y|X)$ and $P(Y|do(X))$. From Fig. 7 we can see that the probabilities estimated by CISNet are more close to the ground-truth, and the probabilities estimated by the baseline model reflect some AU semantic relations such




Figure 6: t-SNE visualization for AU-specific vanilla backbone w/ or w/o CIS module on BP4D dataset.




Figure 7: Differences between $P(Y|X)$ and $P(Y|do(X))$.

as the co-occurrence of AU6 and AU7 which are not suitable for subjects in the samples. Such kind of prior from the training data leads to poor prediction results during inference.

Conclusion

This paper focuses on explaining the why and wherefores of subject variation problem in AU recognition with the help of causal inference theory and providing a solution for subject-invariant facial action unit recognition by deconfounding variable S in the causal diagram via causal intervention. Unlike previous works that made attempt to deal with this problem through subject-specific learning or domain adaptation, we proposed a plug-in causal intervention module named CIS to remove the adverse effect brought by confounder *Subject* in a straightforward way, which could be inserted into almost all frame-based AU recognition model and boost them to a new state-of-the-art. Extensive experiments prove the effectiveness of our CIS module, and vanilla backbones with CIS module inserted achieve state-of-the-art results.

Acknowledgement

This work is in part supported by the PKU-NTU Joint Research Institute (JRI) sponsored by a donation from the Ng Teng Fong Charitable Foundation.

References

- Chen, J.; Liu, X.; Tu, P.; and Aragones, A. 2013. Learning person-specific models for facial expression and action unit recognition. *Pattern Recognition Letters*, 34(15): 1964–1970.
- Chen, Y.; Chen, D.; Wang, Y.; Wang, T.; and Liang, Y. 2021a. CaFGraph: Context-aware Facial Multi-graph Representation for Facial Action Unit Recognition. In *Proceedings of the 29th ACM International Conference on Multimedia*, 1029–1037.
- Chen, Y.; Wu, H.; Wang, T.; Wang, Y.; and Liang, Y. 2021b. Cross-modal Representation Learning For Lightweight and Accurate Facial Action Unit Detection. *IEEE Robotics and Automation Letters*.
- Corneanu, C.; Madadi, M.; and Escalera, S. 2018. Deep structure inference network for facial action unit recognition. In *Proceedings of the European Conference on Computer Vision*, 298–313.
- Cui, Z.; Song, T.; Wang, Y.; and Ji, Q. 2020. Knowledge augmented deep neural networks for joint facial expression and action unit recognition. *Advances in Neural Information Processing Systems*, 33.
- Ekman, P. 1992. An argument for basic emotions. *Cognition and Emotion*.
- Friesen, E.; and Ekman, P. 1978. Facial action coding system: a technique for the measurement of facial movement. *Palo Alto*, 3(2): 5.
- He, K.; Zhang, X.; Ren, S.; and Sun, J. 2016. Deep residual learning for image recognition. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*.
- Hu, X.; Tang, K.; Miao, C.; Hua, X.-S.; and Zhang, H. 2021. Distilling Causal Effect of Data in Class-Incremental Learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 3957–3966.
- King, D. E. 2009. Dlib-ml: A machine learning toolkit. *Journal of Machine Learning Research*, 10(Jul): 1755–1758.
- Li, G.; Zhu, X.; Zeng, Y.; Wang, Q.; and Lin, L. 2019. Semantic relationships guided representation learning for facial action unit recognition. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, 8594–8601.
- Li, W.; Abtahi, F.; and Zhu, Z. 2017. Action unit detection with region adaptation, multi-labeling learning and optimal temporal fusing. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*.
- Li, W.; Abtahi, F.; Zhu, Z.; and Yin, L. 2017. EAC-Net: A region-based deep enhancing and cropping approach for facial action unit detection. In *International Conference on Automatic Face and Gesture Recognition*.
- Mavadati, S. M.; Mahoor, M. H.; Bartlett, K.; Trinh, P.; and Cohn, J. F. 2013. DISFA: A spontaneous facial action intensity database. *IEEE Transactions on Affective Computing*.
- Niu, X.; Han, H.; Yang, S.; Huang, Y.; and Shan, S. 2019. Local relationship learning with person-specific shape regularization for facial action unit detection. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 11917–11926.
- Paszke, A.; Gross, S.; Chintala, S.; Chanan, G.; Yang, E.; DeVito, Z.; Lin, Z.; Desmaison, A.; Antiga, L.; and Lerer, A. 2017. Automatic differentiation in PyTorch. In *NIPS Workshop*.
- Pearl, J.; Glymour, M.; and Jewell, N. P. 2016. *Causal inference in statistics: A primer*. John Wiley & Sons.
- Pearl, J.; et al. 2000. Models, reasoning and inference. *Cambridge, UK: Cambridge University Press*, 19.
- Rubin, D. B. 2005. Causal inference using potential outcomes: Design, modeling, decisions. *Journal of the American Statistical Association*, 100(469): 322–331.
- Shao, Z.; Liu, Z.; Cai, J.; and Ma, L. 2018. Deep adaptive attention for joint facial action unit detection and face alignment. In *Proceedings of the European Conference on Computer Vision*.
- Shao, Z.; Zou, L.; Cai, J.; Wu, Y.; and Ma, L. 2020. Spatio-temporal relation and attention learning for facial action unit detection. *arXiv preprint arXiv:2001.01168*.
- Song, T.; Chen, L.; Zheng, W.; and Ji, Q. 2021a. Uncertain graph neural networks for facial action unit detection. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 1.
- Song, T.; Cui, Z.; Zheng, W.; and Ji, Q. 2021b. Hybrid Message Passing With Performance-Driven Structures for Facial Action Unit Detection. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 6267–6276.
- Tang, K.; Huang, J.; and Zhang, H. 2020. Long-tailed classification by keeping the good and removing the bad momentum causal effect. *arXiv preprint arXiv:2009.12991*.
- van der Maaten, L.; and Hinton, G. 2008. Visualizing Data using t-SNE. *Journal of Machine Learning Research*, 9: 2579–2605.
- Vaswani, A.; Shazeer, N.; Parmar, N.; Uszkoreit, J.; Jones, L.; Gomez, A. N.; Kaiser, L.; and Polosukhin, I. 2017. Attention is All you Need. In *NIPS*.
- Walecki, R.; Rudovic, O.; Pavlovic, V.; and Pantic, M. 2017. Copula ordinal regression framework for joint estimation of facial action unit intensity. *IEEE Transactions on Affective Computing*, 10(3): 297–312.
- Wang, C.; and Wang, S. 2018. Personalized Multiple Facial Action Unit Recognition Through Generative Adversarial Recognition Network. In *ACMMM*.
- Wang, Z.; Li, Y.; Wang, S.; and Ji, Q. 2013. Capturing global semantic relationships for facial action unit recognition. In *Proceedings of the IEEE International Conference on Computer Vision*, 3304–3311.

Xu, K.; Ba, J.; Kiros, R.; Cho, K.; Courville, A.; Salakhudinov, R.; Zemel, R.; and Bengio, Y. 2015. Show, attend and tell: Neural image caption generation with visual attention. In *International conference on machine learning*, 2048–2057. PMLR.

Yang, H.; Yin, L.; Zhou, Y.; and Gu, J. 2021. Exploiting Semantic Embedding and Visual Feature for Facial Action Unit Detection. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 10482–10491.

Yue, Z.; Zhang, H.; Sun, Q.; and Hua, X.-S. 2020. Interventional few-shot learning. *arXiv preprint arXiv:2009.13000*.

Zarins, U. 2018. *Anatomy of Facial Expressions*. Exonicus, Incorporated.

Zen, G.; Porzi, L.; Sangineto, E.; Ricci, E.; and Sebe, N. 2016. Learning personalized models for facial expression analysis and gesture recognition. *IEEE Transactions on Multimedia*, 18(4): 775–788.

Zhang, D.; Zhang, H.; Tang, J.; Hua, X.; and Sun, Q. 2020. Causal intervention for weakly-supervised semantic segmentation. *arXiv preprint arXiv:2009.12547*.

Zhang, X.; Yin, L.; Cohn, J. F.; Canavan, S.; Reale, M.; Horowitz, A.; Liu, P.; and Girard, J. M. 2014. BP4D-spontaneous: A high-resolution spontaneous 3d dynamic facial expression database. *Image and Vision Computing*, 32(10): 692–706.

Zhao, K.; Chu, W.; and Zhang, H. 2016. Deep region and multi-label learning for facial action unit detection. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*.

Zhao, T.; and Wu, X. 2019. Pyramid feature attention network for saliency detection. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 3085–3094.