

The King is Naked: on the Notion of Robustness for Natural Language Processing

Emanuele La Malfa, Marta Kwiatkowska

Department of Computer Science
University of Oxford

{emanuele.lamalfa, marta.kwiatkowska}@cs.ox.ac.uk

Abstract

There is growing evidence that the classical notion of adversarial robustness originally introduced for images has been adopted as a *de facto* standard by a large part of the NLP research community. We show that this notion is problematic in the context of NLP as it considers a narrow spectrum of linguistic phenomena. In this paper, we argue for *semantic robustness*, which is better aligned with the human concept of linguistic fidelity. We characterize *semantic robustness* in terms of biases that it is expected to induce in a model. We study semantic robustness of a range of *vanilla* and robustly trained architectures using a template-based generative test bed. We complement the analysis with empirical evidence that, despite being harder to implement, *semantic robustness* can improve performance on complex linguistic phenomena where models robust in the classical sense fail.

Introduction

In the last decade, deep learning has become the gold standard method to solve complex problems in Natural Language Processing (NLP) (Brown et al. 2020). The range of NLP applications encompasses text classification (Liu et al. 2019), language translation (Liu et al. 2020), and now also ranking systems and large-scale search engines (Wang et al. 2021). With models whose complexity – and consequently, size – has become ‘gargantuan’¹, there is an increasing concern about reproducibility (Liu et al. 2021b) and reliability of those models (Song et al. 2020), as it is known that, even for smaller networks, it is possible to exploit their brittleness with techniques of adversarial machine learning (Zhang et al. 2020). Consequently, concepts of robustness have been transferred from adversarial learning to NLP, resulting in techniques and tools (Li et al. 2020a) that typically check that the network’s decision is invariant to a simple bounded perturbation (word substitution or deletion) for a given input (local robustness), working in the (continuous) embedding space or the (discrete) word neighbourhood. However, NLP still lacks a definition of robustness that properly captures linguistic phenomena and is aligned with human common sense (Xu et al. 2020). There is currently a debate in the NLP community about the internal working of language models,

with some believing they are the ‘foundation’ for the entire discipline (Bommasani et al. 2021) and others arguing that they mostly learn higher-order distributions of words frequency (Sinha et al. 2021).

In this work, we first review the classical notions of robustness adopted in NLP and identify their weaknesses, in terms of the lack of expressiveness and over-reliance on the neural model text representation. Next, to better align the perception of human robustness to that implemented by a neural model, we formalise (local) *semantic robustness* of NLP as a notion that generalizes local discrete robustness through measuring robustness to linguistic rules, rather than to word substitution or deletion. This allows us to define (global) semantic robustness for a linguistic task such as sentiment analysis, which can be extended to higher-order tasks. We contribute to the debate in the NLP community by performing a systematic comparison, complemented by an evaluation of different architectures, of the classical notions of robustness in NLP. We further show that with *semantic robustness* we can evaluate the performance of a model on cogent linguistic phenomena, which are of interest for both the NLP and the linguistics community. We achieve this by proposing an assessment framework and a simple, yet effective, test bed based on data augmentation. Last but not least, we wish to highlight the issue of NLP robustness, which for the last few years has over-focused on trivial and often machine-centric symbol manipulation. Using the terminology drawn from cyber-security, this work is a ‘purple-team’ effort to align the key performance indicators of the ‘red-team’ – whose role is to exploit NLP models with any kind of vulnerability – with those of the ‘blue-team’, a.k.a. the defenders, who aim to adopt a semantic notion of robustness that implies robustness to linguistic phenomena.

Related Work

Brittleness of neural network models is a serious concern, both theoretically (Biggio et al. 2013; Szegedy et al. 2014) and practically, including Natural Language Processing (NLP) (Belinkov and Bisk 2018; Ettinger et al. 2017; Gao et al. 2018; Jia and Liang 2017; Liang et al. 2017; Zhang et al. 2020) and more recently complex Masked Language Models (MLM) (Li et al. 2020b; Sun et al. 2020). In NLP, attacks are usually conducted either at character or word level (Ebrahimi et al. 2017; Cheng et al. 2018), or

Copyright © 2022, Association for the Advancement of Artificial Intelligence (www.aaai.org). All rights reserved.

¹GPT-3’s full version has 175 billion learning parameters.

at the embedding level, exploiting (partially or fully) vulnerabilities in the symbols’ representation (Alzantot et al. 2018; La Malfa et al. 2021). Brittleness of NLP does not pertain only to text manipulation, but also includes attacks and complementary robustness for ranking systems (Goren et al. 2018). Neural network robustness naturally complements the perspective offered by brittleness as it involves the certification of a model against a wide range of attacks (Huang et al. 2017). In NLP, similarly to computer vision (Akhtar and Mian 2018), the majority of works have adopted the narrow notion of robustness, in terms of invariance to minor perturbations of an input text (Gowal et al. 2018; Jia et al. 2019; Dong et al. 2021; La Malfa et al. 2020), while only a minority have contested these limitations, either implicitly (Ribeiro et al. 2020) or explicitly (Morris 2020; Morris et al. 2020a; Xu et al. 2020), mainly due to the difficulty of automatically generating semantically involved test beds (Feng et al. 2021). Although adversarial data augmentation in NLP is well established (Morris et al. 2020b), robustness to semantically coherent, yet possibly diverging, examples is still in its ‘adolescence’ (Ribeiro, Singh, and Guestrin 2018), as many highly accurate NLP models cannot recognize cogent linguistic phenomena even on low-order tasks such as binary classification (Barnes, Øvrelid, and Veldal 2019).

NLP Robustness: a Tale of Two Perspectives

In this section we discuss the merits of local robustness in NLP, analysing existing concepts and highlighting the issues with local continuous robustness. We then introduce the notion of semantic robustness, aimed to better align the perception of human robustness to those implemented by neural models, together with an assessment framework. In the next section, we complement the methodology part with an experimental evaluation, where neural networks’ robustness is tested against linguistic phenomena. We complete the paper with a study of the inductive biases that different notions of robustness are expected to induce in a trained model.

Notation. We will refer to $f(\cdot)$ as a generic neural network that solves a task T for an instance s , which is represented as a piece of text written in natural language (e.g., sentiment analysis). W.l.o.g., we will assume that texts are represented as lists of words (features), namely $s = (v_1, \dots, v_l)$. We will denote with $x \in \mathbb{R}^{ld}$ a text s whose l features have been mapped to vectors of real numbers through an embedding, $\mathcal{E} : \mathcal{V} \rightarrow \mathbb{R}^d$, where \mathcal{V} is a finite vocabulary of words. We refer to a component of x along a generic embedding axis as $x^{(i)} \in \mathbb{R}$, $i \in \{1, \dots, d\}$. Since the majority of the embedding spaces are injective non-surjective functions, the notation $x_v = (x_{v_1}, \dots, x_{v_l})$ will serve to denote an embedded text x that further admits, for each vector x_{v_i} , a preimage in the vocabulary space, i.e., $\forall x_{v_i} \in x_v \exists! v \in \mathcal{V}. \mathcal{E}(v) = x_{v_i}$. With a slight abuse of notation, we will denote with $\mathcal{E}(s) = (x_{v_1}, \dots, x_{v_l})$ a text whose words have each been embedded through \mathcal{E} . Finally, we will assume that the first operation of the model $f(\cdot)$ involves a transformation through an embedding representation \mathcal{E} .

Classical Notions of NLP Robustness

We begin by discussing the concept of local continuous robustness, which is widely used in computer vision and has been applied to NLP (Huang et al. 2019a; Jia and Liang 2017; La Malfa et al. 2020). We then consider local discrete robustness, which manipulates symbols rather than embedding vectors (Alzantot et al. 2018). We show that the former notion can be reduced to the latter: nonetheless, both definitions only allow one to express robustness to a limited number of linguistic phenomena. We extensively discuss the advantages and drawbacks of those two notions.

Definition 1 (Local Continuous Robustness). A model $f(\cdot)$ is locally robust to ϵ -bounded perturbations when, given a task T and one of its instances x , it holds that $\forall x' \in Ball_\epsilon(x), f(x) = f(x')$, where $Ball_\epsilon(x) = \{x' . \|x - x'\|_p \leq \epsilon\}$, $\| \cdot \|_p$ is an L_p norm of choice and $\epsilon \geq 0$ a (small) real number.

Observation 1. Natural language is discrete while local continuous robustness is defined over a dense representation. Standard embedding techniques (Mikolov et al. 2013; Pennington, Socher, and Manning 2020) define the word-to-vector mapping over a continuous space, with the input vocabulary discrete and finite (e.g., characters, words, sentences) and the output dense and uncountable. On the other hand, natural language is discrete and allows for finite, yet combinatorial, outcomes. In this *hybrid* setting, ϵ -bounded robustness implies that any vector in this dense ϵ -bounded region is safe. This assumption is linguistically inconsistent, as a network may present a decision boundary where an adversarial attack that is not a proper word limits the verification or severely reduces the safe region. We illustrate this issue in Figure 1.

Definition 2 (Local Discrete Robustness). A model $f(\cdot)$ is locally robust to discrete perturbations when, given a task T and an instance x_v embedded from s , it holds that $\forall x'_v \in D-Ball_\epsilon(x_v), f(x_v) = f(x'_v)$, where $D-Ball_\epsilon(x_v) = \mathcal{E}(\mathcal{V})^l \cap Ball_\epsilon(x_v)$.

We exemplify the differences between $Ball_\epsilon$ and the corresponding $D-Ball_\epsilon$ in Figure 2.

Proposition 1. Local continuous robustness implies local discrete robustness, but the converse is generally false.

Proof. From a mathematical perspective, $D-Ball_\epsilon(x) \subseteq Ball_\epsilon(x)$ but the opposite is not true. For $\epsilon = 0$, both $D-Ball_\epsilon$ and $Ball_\epsilon$ are singletons. ■

Observation 2. Continuous and discrete robustness allow but limited syntax manipulations. As demonstrated empirically by many existing works in the literature (Alzantot et al. 2018; Jia et al. 2019; Huang et al. 2019a; Dong et al. 2021), both formulations of robustness only allow for robustness testing against symbol-to-symbol substitutions or deletions. The limited degree of freedom of an operator that locally substitutes a word with other words makes it hard, if not impossible, to test for robustness against paraphrases. As an example, if a model $f(\cdot)$ is robust for the sentence “*the movie was good*”, which implies correct classification for

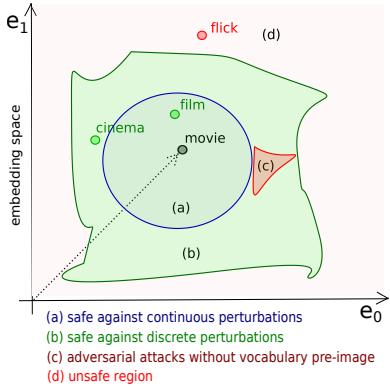


Figure 1: In general, local continuous robustness is an ill-posed property for NLP. A model can be robust to a large surface of attacks in the input neighbourhood (green patch (b)), yet a small region of adversarial attacks (red patch (c)) invalidates the verification of larger regions. In the example, the safe input neighbourhood (blue patch (a)), a convex region that includes safe replacements, cannot grow any further without violating robustness by encroaching on patch (c). Non-convex representations for an input neighborhood (patch (a)) are possible, but computationally expensive and not used in practice.

the texts “*the film was good*”, “*the film was nice*”, etc., we cannot say the same for the sentence “*an enjoyable thriller*”. From the linguistic perspective, this problem arises since the frequency of words in natural language follows the Zipf’s law (Zipf 2013), where rare terms and constructs – hence edge cases – occur more frequently than in other natural phenomena.

Observation 3. There is no guarantee that perturbations in both discrete and continuous settings do not violate the task under consideration. As the methods that implement both discrete and continuous robustness allow for weak supervision in the choice of the replacements, a perturbation can diverge from the task under consideration. It is well known that many recent embeddings have been developed to be faithful to a (static) version of the “distributional hypothesis” (Baroni, Dinu, and Kruszewski 2014), and thus it is not unusual to find words like “*bad*” and “*good*” close to each other in the representations. This could lead to potentially disastrous effects when balancing local robustness, e.g., (Gowal et al. 2018), with accuracy, especially for low-order tasks such as sentiment analysis.

A Semantic Notion of Robustness

We now introduce a notion of robustness that goes beyond word replacements, and thus permits an assessment of the brittleness to linguistic phenomena that are cogent to humans. To do so, we first need to introduce some notation.

Definition 3 (Oracle). An Oracle Ω solves a task T for any input s that is compliant with T , while it rejects all those instances that are not. We denote with $\Omega \models s$ the act of

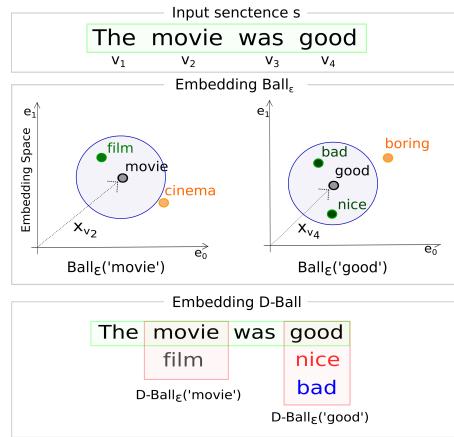


Figure 2: $Ball_\epsilon$ (top) and $D\text{-}Ball_\epsilon$ (bottom) representations of two words from the input sentence “*the movie was good*” (s in our notation). For the same value of ϵ , $Ball_\epsilon$ contains all the discrete replacements of the equivalent $D\text{-}Ball_\epsilon$ plus all the vectors (infinitely many) that cannot be mapped back to the vocabulary \mathcal{V} (inside each blue ball around an input word).

solving a task, and with $\Omega \not\models s$ the rejection. Solving and rejection are mutually exclusive.

Observation 4. An Oracle is an augmented, idealized linguistic model. There are two cogent differences between an Oracle Ω and a standard model $f(\cdot)$: (i) a model can be wrong on samples from T (e.g., misclassifications) while the Oracle is always right about its decisions; (ii) the Oracle (certainly) rejects inputs that are not compliant with the task T .

Example 1. An Oracle for sentiment analysis. Given a sentiment analysis task T for movie reviews, an Oracle Ω correctly classifies any text that expresses a judgment about the movie. As an example, “*the movie was (not) good*” will be classified as positive (negative). An Oracle alternatively rejects any piece of text that is inconsistent with T , i.e., all those texts that do not explicitly (or implicitly) express a judgment about a movie. An example is the text “*recipe of risotto with mushrooms: [...]*”, which is rejected. From a practical perspective, a classifier $f(\cdot)$ admits misclassifications in the sense that its accuracy may not be maximal (i.e., less than 1.) and it further cannot reject inputs that are not compliant with the task: in the case it does, the task is not fulfilled perfectly.

Definition 4 (Linguistic Rule). A linguistic rule is a symbolic function that manipulates a text s according to a linguistic phenomenon and a task T , whose generated texts S' , along with the original input, are not rejected by Ω . Formally, $R : (s, T) \mapsto S' . \forall s' \in S', \Omega \models s \wedge s'$.

Observation 5. Linguistic rules are flexible symbolic methods. Since a linguistic variation of an input can be very different from the original text, a rule should be allowed to add/remove/replace words while remaining com-

pliant with the task T . As an example of a simple rule, one can think of *verb negation* that acts on a text and negates the action expressed by the subject (if any). While this task is often trivial for humans, fully algorithmic solutions to this problem are still limited in their capabilities (Guo et al. 2018). *Hybrid* methods, based on synthetic data augmentation, humans-in-the-loop and deep MLM (Feng et al. 2021; Lin et al. 2019; Huang et al. 2019b), constitute currently an active area of NLP research. One viable way to generate the replacements is to use template-based data-augmentation techniques (as employed in this paper and detailed in Experimental Evaluation). More complex approaches involve MLM with humans-in-the-loop who validate the generated perturbations. While for the generative process the MLM can be trained to be controlled through textual ‘seeds’ (e.g., in the spirit of the works by (Wu et al. 2021) or (Madaan et al. 2020)), humans play the role of the Oracle.

Example 2. A rule for *shallow negation*. For a sentiment analysis task \bar{T} with positive and negative instances and a positive instance s “*the movie was good*”, the *shallow negation* rule R negates the sentiment expressed by s , and hence valid perturbations generated by R on s are “*the movie was not good*”, “*a bad film*”, but also more involved examples like “*it is false that the movie is good*”, etc. We name this rule shallow negation as it does not allow for nested negations, regardless of their grammatical consistency (i.e., “*it is false that the movie wasn’t good*” cannot be generated by R on s).

Definition 5 (Local Semantic Robustness). Formally, given a model $f(\cdot)$, a linguistic rule R , an input s from a task T , a measure of the performance of $f(\cdot)$ on T , namely $p \in [0, 1]$ (with 0. denoting a random guess and 1. perfect accuracy), a small number $\tau \geq 0$, and a measure of performance p' on samples S' generated by $R(s, T)$, we say $f(\cdot)$ is τ -semantically robust for R and s if it holds that $\mathbb{E}_{s' \sim S'}[p'] \geq p - \tau$, with $\min(0, p - \tau) = 0$.

We further say that a model is *bounded invariant* to a rule R when it holds that $p - \tau \leq \mathbb{E}_{s' \sim S'}[p'] \leq p + \tau$.

Informally, a model $f(\cdot)$ that correctly classifies an instance s of a task T is semantically robust to a linguistic rule R when it exhibits at least the same performance on the set S' of perturbations s' generated by applying R to s . We further observe that this formulation allows for the performance p' to even surpass p , so this notion entails that $f(\cdot)$ is no worse at correctly solving T for S' than it is at solving any other task, and is hence a stronger notion than *bounded invariance*.

Observation 6. Local semantic robustness is linguistically meaningful. The notion is local as linguistic rules act on a single text s . It is further inherently linguistic as the transformation R of an input text acts at the syntax level but then the Oracle’s reject phase guarantees it has preserved the semantics of each s' w.r.t. T .

Observation 7. Local semantic robustness is entailed by linguistic generalization, but not the other way round. Linguistic robustness is different from generalization on unseen

test cases. The former is entailed by the latter, while the other way round is not necessarily true. Semantic robustness is defined over a rule while generalization is a more general and hard to obtain/optimize objective.

Proposition 2. Local semantic robustness can be reduced to local discrete robustness, but not to local continuous robustness.

Proof. For local discrete robustness, it is straightforward to define a rule that generates perturbations according to the definition of local discrete robustness. In this sense the semantic rule R involves extracting the replacements in the embedding’s neighborhood of each input word.

As regards local continuous robustness, the invariance over all the input texts s' in an ϵ -ball cannot be mapped back to the embedding (a.k.a. input) vocabulary \mathcal{V} by any combination of linguistic rules as they act, by definition, at the symbol level. Since the majority of continuous embeddings are injective non-surjective functions, almost all the vectors in any non-empty region of the space cannot be mapped back to a proper entry of \mathcal{V} . ■

Definition 6 (Semantic Robustness). This notion extends local semantic robustness beyond the single instance and to a specific task T . A model $f(\cdot)$ exhibits global semantic robustness (or in general semantic robustness) to a rule R and a task T when it is locally semantically robust for any input s' generated by applying R to a test set.

Assessment Framework for Semantic Robustness A sufficient condition for quantifying the semantic robustness of a model on an NLP task is that it is possible to measure the performance of such a model on unseen input texts. In this sense, we can measure the semantic robustness of a model $f(\cdot)$ that solves a task T by comparing its performance p with the performance p' of the model on an unseen test bed that contains one or more semantic phenomena.

We now describe some illustrative examples of measuring semantic robustness, firstly for sentiment analysis and then for more involved NLP tasks.

Example 3. Robustness to *shallow negation* in sentiment analysis. Given a sentiment analysis task with positive and negative instances, a model $f(\cdot)$ trained on a dataset (S, Y) and validated on (S_{test}, Y_{test}) is robust to shallow negation when $\forall s \in S_{test}, \forall s' \in S' = R(s, T), (\Omega \models s \wedge s') \Rightarrow \mathbb{E}_{s' \sim S'}[p'] \geq p - \tau$ for some $\tau \geq 0$, with R the negation rule that acts on a specific text and negates the sentiment expressed by s . In this sense, p represents the accuracy of the trained model on (S_{test}, Y_{test}) , while p' is the accuracy measured on a subset of samples that contain specific linguistic phenomena. We remark that a test bed can be handcrafted, as we show in our paper, or distilled from existing datasets, as described in (Barnes, Øvreliid, and Velldal 2019).

Example 4. Semantic robustness in high-order NLP tasks. We now briefly sketch how we would approach the measurement of semantic robustness for higher-order NLP tasks. For Question and Answer (QA) tasks, a measure of robustness can be quantified as the gap between the ‘unexpectedness’ of an Answer when the Question does/doesn’t contain a lin-

guistic phenomenon. In Natural Language Inference (NLI), directly applying our framework would be straightforward since NLI is reducible to a classification task. In the same way, when Read and Comprehension (RC) is pursued in the form of a classification task, the evaluation of semantic robustness would be similar to sentiment analysis or NLI, whereas when the answer requires re-elaborating the input, the measurement of semantic robustness would be similar to QA (with possibly a different evaluation metric for T).

Experimental Evaluation

We next conduct an extensive experimental evaluation² designed to answer the following research questions: (i) whether models robust in the classical sense are also semantically robust; (ii) whether robustness to specific linguistic phenomena is a by-product of training accurate NLP classifiers; (iii) whether, for different architectures, augmented supervised training – with texts that contain a specific linguistic phenomenon – induces generalization on unseen test samples that contain the same phenomenon; (iv) whether it is possible to train models that are both accurate and semantically robust, and (v) to what extent unsupervised learning contributes to semantic robustness.

We conduct the experiments on models trained – or fine-tuned through data augmentation – on the Stanford Sentiment Treebank dataset (SST-2) (Socher et al. 2013) and on the dataset collected by (Barnes, Øvrelid, and Velldal 2019). The advantages of this approach are two-fold. Firstly, human experts have collected/handcrafted sentences whose syntax/semantics is rich and the level of noise restrained. Secondly, since in NLP spurious patterns and over-fitting play a crucial role during training whose influence is hard to estimate and quantify, cogent compactness of those datasets makes it relatively easy to assess the results. To further estimate the robustness on linguistic phenomena, in the spirit of the evaluation done in (Huang et al. 2019b), we utilise a template-based method, whose details are given below, for generating augmented samples for a selection of linguistic phenomena to create a test bed, which we use for systematic evaluation of semantic robustness. In order to validate the soundness of our generative test bed, we compare the performance of our rule-generated semantically robust models from our benchmark to those examples in (Barnes, Øvrelid, and Velldal 2019) that exhibit the same linguistic phenomenon, showing comparable accuracy.

Linguistic phenomena. Following the work in (Barnes, Øvrelid, and Velldal 2019), we have chosen interesting linguistic and para-linguistic phenomena, taking care to exclude those that require external knowledge to be solved

²The code for full reproducibility of the experiments is available at <https://github.com/EmanueleLM/the-king-is-naked>. All of the experiments have been conducted on a Fedora 32 mid-end laptop equipped with 16GB of RAM and an Intel-i5 CORE 8th-generation. All the neural network models have been built, trained and tested with Keras (Chollet et al. 2015), while for experiments that involved BERT (Devlin et al. 2018) we relied on the PyTorch implementation (Paszke et al. 2019).

(i.e., not explicitly expressed in the sentence). As an example, consider the review “*This movie is another Vietnam*”, which can be correctly classified as negative if the model has some knowledge of that specific way of saying (i.e., exogenous knowledge). We now briefly describe the linguistic phenomena that are the object of our robustness evaluation:

- *Shallow negation*: when the sentiment of a sentence is negated. We do not consider nested negations, which make the recognition of the phenomenon considerably harder (Wiegand et al. 2010; Socher et al. 2013; Pröllochs, Feuerriegel, and Neumann 2015).
- *Mixed sentiment*: when phrases of different polarity appear in the same sentence (Kenyon-Dean et al. 2018; Barnes, Øvrelid, and Velldal 2019). We only consider texts where the overall sentiment is still not ambiguous for a human.
- *Irony/sarcasm*: when a sentence makes some premises that are then violated (Hao and Veale 2010). This is known to be one of the hardest, yet pervasive, linguistic phenomena of human language.

Template-based linguistic rules. In addition to the test beds provided by (Barnes, Øvrelid, and Velldal 2019), in our work we consider a template-based method for generating augmented samples that contain a specific linguistic phenomenon. We pre-define a selection of templates for which we know the corresponding output labels (i.e., *positive* or *negative*). In a template, part of the text is fixed while the remaining part is symbolically represented by tokens which are iteratively replaced by combinations of words from candidate perturbation sets. The augmentation preserves the semantics of the sentence while introducing a linguistic phenomenon (such as *shallow negation*). In our implementation of the rules, a perturbation cannot change the template’s label: in this sense, the rejection phase (see Definition 3) is embedded in the generative pipeline, while a process that involves an MLM and generations that are possibly label-changing might be supervised by a human. Examples of templates for each linguistic rule are included in Table 3, along with candidate replacements for each token in Table 2.

Comparative Study

We compare architecturally different models on the three linguistic phenomena we previously introduced. We conduct an extensive evaluation on four neural architectures, namely fully connected (FC), convolutional (CNN) (Zhang, Zhao, and LeCun 2015), Long Short-Term Memory (LSTM) (Hochreiter and Schmidhuber 1997) and self-attention (Vaswani et al. 2017). We choose the number of hidden units of each layer so that the number of parameters is approximately the same and in the order of 40K³. For each

³Each input text is 25 words long (eventually padded or cut), while each word is mapped to a vector of real numbers through a 50-dimensional embedding, pre-trained on the SST-2 task (Chollet et al. 2015). Each network is composed of 3 layers, where the topology of the last two is shared, i.e., respectively a 32 hidden units ReLU and a 2 hidden units softmax layer (both are dense). The first layer depends on the specific topology under examination.

	Train	FCs	CNNs	LSTMs	Self-attention
Shallow Negation	<i>Vanilla</i>	0.4034 ± 0.0214	0.4032 ± 0.0124	0.4771 ± 0.0143	0.4790 ± 0.0059
	<i>Augmented</i>	0.4062 ± 0.0167	0.4249 ± 0.0255	$0.6387 \pm 0.0387^*$	$0.5954 \pm 0.0027^*$
Mixed Sentiment	<i>Vanilla</i>	0.4707 ± 0.0360	0.4986 ± 0.0415	0.5110 ± 0.0251	0.5487 ± 0.0099
	<i>Augmented</i>	0.4912 ± 0.0339	0.5271 ± 0.0387	$0.6357 \pm 0.0317^*$	0.5617 ± 0.0048
Sarcasm	<i>Vanilla</i>	0.5136 ± 0.0504	0.4681 ± 0.0327	$0.5578 \pm 0.0128^*$	0.5240 ± 0.0132
	<i>Augmented</i>	0.5297 ± 0.0657	0.4678 ± 0.0317	0.4807 ± 0.0197	$0.6236 \pm 0.0218^*$

Table 1: Comparison of accuracy of 20 *vanilla* and *augmented* models obtained for four different architectures (FCs, CNNs, LSTMs and self-attention), on three linguistic phenomena (*shallow negation*, *mixed sentiment* and *sarcasm*). All the networks have been trained on the SST-2 dataset. *Augmented* models are *vanilla* models fine-tuned on the linguistic rules of interest. Symbol $*$, when present, means that the improved performance (from *vanilla* to *augmented*, or the other way round) is statistically significant. Interestingly, *sarcasm* is harder to learn and models fine-tuned on this phenomenon perform as well as their *vanilla* counterparts (when not worse).

Tokens	Replacements
@NEGATIVE@	'bad', 'poor', 'boring', [...]
@POSITIVE@	'good', 'nice', 'fantastic', [...]
@NAME@	'Uma', 'Bruce', 'Sandra', [...]
@SURNAME@	'Thurman', 'Willis', 'Bullock', [...]
@CATEGORY@	'thriller', 'horror', 'comedy', [...]
@BOOLFALSE@	'false', 'wrong', 'incorrect', [...]
@AUGMENT@	'very', 'extremely', 'incredibly', [...]

Table 2: Candidate perturbation sets used to generate combinations of replacements in template-based texts (Table 3).

linguistic phenomenon, we analyse and compare the robustness of 20 models trained on *plain* SST-2 dataset (i.e., no semantic data augmentation of any kind) and then on a *semantically augmented* version of the same dataset (details of the augmentation are provided in the relevant subsection).

Vanilla Models. For each linguistic phenomenon introduced in the previous section, we analyse and compare the robustness of 20 models trained on the SST-2 dataset without augmentation. For FCs and CNNs the average accuracy on the SST-2 test set is 0.8993 ± 0.0029 and 0.9077 ± 0.0038 respectively, while the accuracies of LSTMs and self-attention are 0.9101 ± 0.0033 and 0.8963 ± 0.0015 . We report in Table 1 the results of each population for the three linguistic phenomena in this study. Self-attention and RNN-LSTMs are the best performers, while FCs and CNNs have lower accuracy in all the three tasks. Interestingly, none of the models, despite having a high accuracy on the test set, is able to recognize any linguistic construct we tested. On the one hand, this analysis, which can guide the design when seeking to enforce appropriate inductive biases of a neural architecture (Kharitonov and Chaabouni 2020), provides additional evidence for the vast literature on the limitations of accuracy when judging the linguistic performance of an NLP model (Socher et al. 2013; Kenyon-Dean et al. 2018; Barnes,

tion (e.g., self-attention will have a self-attention layer, LSTM a Long Short-term Memory cell, etc.): the first layer has 32 hidden units for the FCs, 44 ReLU kernels of size 3 for the CNNs, 75 tanh hidden units for the LSTMs and 32 ReLU hidden units for the self-attention networks.

Velldal, and Øvreliid 2021). On the other hand, it motivates our next step, which involves fine-tuning the same architectures on texts that exhibit these linguistic phenomena.

Semantic robustness through data augmentation. In this section, we study how – and to what extent – data augmentation, along with architectural inductive biases, can be used to inject semantic robustness to different linguistic phenomena. We re-trained the models of the previous section by adding samples from (Barnes, Øvreliid, and Velldal 2019) that contain one of the specific rules used previously to the training set, up to a multiplicative factor to balance the large number of samples of the SST-2 dataset⁴. While for a multiplicative factor of 500 none of the models exhibit any improvement in the semantic tasks, for a multiplicative factor of 750 we observe some improvement in LSTMs and self-attention. While the experiments suggest that FCs and CNNs cannot learn any of the three linguistic phenomena we studied, LSTMs and self-attention networks benefit from data augmentation. With reference to Table 1, both LSTMs and self-attention improve considerably on *shallow negation*. On *mixed sentiment*, augmented LSTMs substantially improve over the *vanilla* counterpart, while self-attention does not seem to exploit the additional information (despite a slight improvement over the *vanilla* case). Finally, data augmentation allows self-attention to improve significantly on *sarcasm*, though the same regime is detrimental for LSTMs, where the *vanilla* networks consistently outperform those trained on augmented data. Finally, for a multiplicative factor of 1K or superior, we observe a detrimental effect on the robustness of each model that is comparable to the *vanilla* SST-2 training.

Classic Robustness is Linguistically Brittle

We have compared *robust* models trained with IBP (Interval Bound Propagation) (Gowal et al. 2018) with their *vanilla* counterparts. For different values of $\epsilon = (0.001, 0.01)$ in the L_∞ -norm, which makes the model-to-model results easy to compare (La Malfa et al. 2020), and an embedding diameter

⁴With the SST-2 train set that accounts for approximately 112K input samples and each semantic rule that generates roughly 500 – 1000 new samples, semantic data augmentation with a multiplicative factor of 1 accounts for additionally 1K samples, etc.

	<i>Shallow Negation</i>	<i>Label</i>
'This @CATEGORY@ movie is not @AUGMENT@ @NEGATIVE@.'		<i>positive</i>
'It is @BOOLFALSE@ that this @CATEGORY@ movie is @AUGMENT@ @POSITIVE@.'		<i>negative</i>
	<i>Mixed Sentiment</i>	
'Despite @NAME@ @SURNAME@ acted well, this @category@ movie is @augment@ @negative@.'		<i>negative</i>
'A @AUGMENT@ @NEGATIVE@ plot for a @AUGMENT@ @POSITIVE@ movie.'		<i>positive</i>
	<i>Sarcasm</i>	
'Starring @NAME@ @SURNAME@ i would prefer to be killed rather than watching this @CATEGORY@ movie.'		<i>negative</i>
'Please throw this @AUGMENT@ long @CATEGORY@ movie into the ocean, and thank me later.'		<i>negative</i>

Table 3: Examples of template-based reviews, along with the ground truth label, used to generate sentences that contain the linguistic phenomena studied in the paper.

	Accuracy (Barnes et al., 2019)	Accuracy (Our Benchmark)
Shallow Negation	0.8552	0.7928
Mixed Sentiment	0.6024	0.6974
Sarcasm	0.7111	0.8455

Table 4: Summary of BERT semantic robustness on different linguistic phenomena, tested on samples from (Barnes, Øvrelid, and Veldal 2019) (left column) and from our template-based benchmark (right column). For these results, a BERT model has been fine-tuned on the SST-2 dataset.

of approximately 3.17, we assess IBP-induced robustness on semantic rules. Interestingly, their performance is comparable (when not worse) to the brittle counterparts for all the linguistic phenomena we analyse, thus validating our previous Observation 2, i.e., that models robust in the classical sense have an extremely limited syntax/semantic manipulation capability. Results are reported in Table 5.

Accuracy is a red herring: the BERT case

We analyse the relationship between semantic robustness and accuracy of a Masked Language Model (MLM): while it is known that MLMs have an improved accuracy on out-of-distribution (OOD) data (Hendrycks et al. 2020), there is no clear agreement on the nature of the semantic phenomena, i.e., whether they are *linguistic outliers* or OODs. Although in deep learning a trade off has been observed between the classical notions of robustness and accuracy (Tsipras et al. 2018), semantic robustness does not seem to exacerbate this phenomenon. We fine-tuned the BERT language model (Devlin et al. 2018) on the SST-2 dataset and tested its robustness on the linguistic phenomena we introduced in the previous section.

Despite an accuracy of 0.90, which is in line with the accuracy of the (simpler) architectures we tested previously, BERT’s semantic robustness is considerably higher than the “shallow” counterparts (BERT has 16 hidden layers, the models in our benchmark 3). BERT has an accuracy of 0.7928 on *shallow negation*, 0.6974 on *mixed sentiment*, and 0.8445 on *sarcasm*⁵. The linguistic phenomenon where

⁵Due to the high computational cost of fine-tuning BERT, we could not carry out an extensive evaluation correlated by an std interval, as done for the simpler networks.

BERT performs worst is *mixed sentiment*, as: (i) a few recent works point out the limitations of MLM models such as BERT when learning complex syntactic/semantic constructs (Sinha et al. 2021); (ii) we have shown in our previous evaluation that self-attention (along with any other model) is especially brittle to that linguistic construct, despite the layer’s name suggesting the opposite. In general, we interpret this linguistic performance as a result of the huge amount of unsupervised training (i.e., the masked language prediction) to which BERT is subjected before being fine-tuned on our supervised task: in this sense, the phase of pre-training, which shapes the dynamics of BERT’s contextual embeddings, enables it to considerably outperform shallow models on the linguistic phenomena.

We finally validate the results of (Barnes, Øvrelid, and Veldal 2019), proving that on their challenging dataset, which contains texts from other non-movie-review datasets (so certainly out of distribution samples), BERT has an accuracy of 0.8552, 0.6024 and 0.7111 on respectively *shallow negation*, *mixed sentiment* and *sarcasm*. This therefore justifies that the task that we set up with our synthetic augmentation through templates is a solid alternative benchmark for semantic robustness. We summarize the results in Table 4.

Ablation Study of BERT. We performed an ablation study of BERT to assess the role of the stacked embeddings to semantic robustness. We hence trained different semantic classifiers on top of a decreasing number of BERT embedding layers. We then measured the semantic robustness on *shallow negation*, *mixed sentiment* and *sarcasm* on samples from (Barnes, Øvrelid, and Veldal 2019): we found that, despite the accuracy on the task (SST-2) being strongly correlated with the depth of the BERT embedding, semantic robustness is not, as depicted in Figure 3. While the best performing layer is the penultimate? (an interesting phenomenon that is already known in the literature (Rogers, Kovaleva, and Rumshisky 2020)), we could not find a layer that performed the best on all the tasks, a result that leads us to conclude that stacked attention embeddings are fundamental but their internal representation w.r.t. linguistic phenomena (i.e., the ‘semantics of BERT’) is still poorly understood. To complement the analysis, we tried to disentangle the role of pre-training from that of the embedding depth and attention (which are considered in the design of each BERT hidden layer) by training a very deep LSTM, with 100 input words

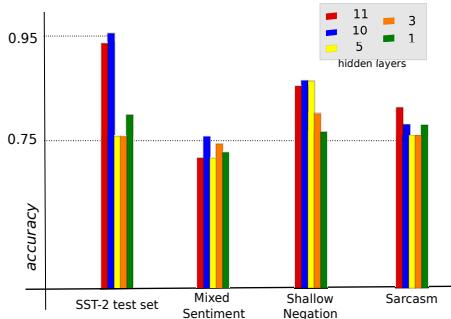


Figure 3: Ablation study of BERT on (Barnes, Øvrelid, and Velldal 2019), measuring accuracy for 5 different network depths. While depth plays a fundamental role in achieving accuracy on a test set (*SST-2*), and certainly plays a role (albeit minor) on *shallow negation*, it seems not to be correlated to the model performance on *mixed sentiment* and *sarcasm*.

and an embedding size of 100, which we then tested on the same semantic phenomena as in the previous evaluation. Interestingly, despite an accuracy of 0.9 on the *SST-2* test set, the accuracy on *shallow negation* is 0.5789, 0.6684 on *mixed sentiment* and 0.7 on *sarcasm*. Although we cannot conclude anything definite, we suspect that the role played by massive pre-training (next word/sentence prediction) is much more important than that of depth and attention, which is in agreement with observations emerging from other recent studies (Liu et al. 2021a).

Robustness Induced Biases

In this section we examine the relationship between common inductive biases that have inspired the design of machine learning algorithms for the past decades (Mitchell 1980), and recently also neural networks (Kharitonov and Chaabouni 2020), connecting them to the notions of robustness we dissected in the previous section. In particular, we compare local continuous to local semantic robustness.

Minimum Cross-validation Error. There is empirical evidence in the literature (Huang et al. 2019a; Jia et al. 2019) that continuous robustness does not naturally induce better performance on trained models. Indeed, most of the models that are trained to be robust are less accurate than the brittle counterparts. This side-effect is caused by the margin that is propagated through the network to the output to induce invariance to nearest neighbours of a given input. “Shielding” the model with a thick margin of possibly unrelated terms leads to an inconsistent treatment of different sentences (as noted in Observation 2, human language abounds in edge cases). This is testified by further experiments shown in Figure 4 (top). Concerning semantic robustness, generalization on cogent linguistic rules does not necessarily benefit a model’s performance, as demonstrated by experiments we conducted on 30 networks trained to be semantically robust against *shallow negation* vs. their vanilla counterpart: both populations have been trained on the *SST-2* dataset (Socher et al. 2013). Robustness is enabled through simple data aug-

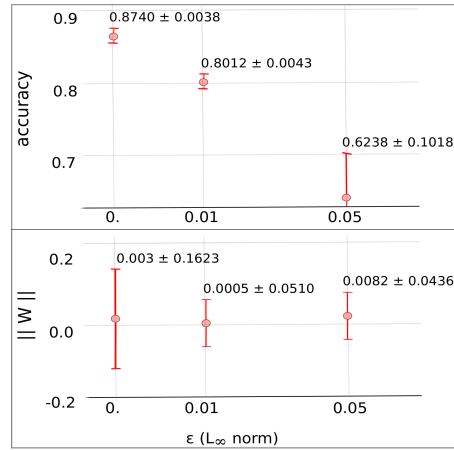


Figure 4: On the top plot, we show average accuracy of 30 trained FC models on the *SST-2* dataset, compared for different values of ϵ -robustness (w.r.t. the L_∞ -norm). For ϵ equal to 0., a model is not robustly trained, and otherwise it is through IBP (Gowal et al. 2018). There is clear a trade-off between robustness and accuracy. On the bottom plot, the average norm of the models’ parameters indicates that robust models tend to have lower variance, and hence arguably lower complexity.

mentation on the dataset provided by Barnes et al. (Barnes, Øvrelid, and Velldal 2019), whereas the test is performed on unseen sentences that exhibit the same linguistic phenomenon. While the vanilla networks have an average accuracy of 0.9036 ± 0.0019 on the test set and 0.4916 ± 0.0074 on the *shallow negation* test set, those that have been robustly trained have an accuracy of 0.8838 ± 0.0049 and 0.5491 ± 0.0124 , respectively.

Minimum Description Length. Local continuous robustness is known to be a strong regularizer (Gowal et al. 2018). In fact, classical methods used to induce local robustness for NLP (such as IBP), which propagate through all the embedding dimensions and thus amplify the noise, nonetheless play an important role as they smooth out the network’s hidden activations. We report the results of experiments that we conducted that support this hypothesis in Figure 4 (bottom). As regards semantic robustness, we cannot conclude anything definitive but the evidence suggests that semantically robust models are not necessarily smoother than the vanilla counterparts. We compared the weights’ norm of 30 networks trained to be robust against *shallow negation* vs. their vanilla counterparts (see previous paragraph for details). While the difference between the performance of the two networks on unseen texts that contain that linguistic phenomenon is substantial, there is very little difference in the norm of the two populations, which are respectively 0.0017 ± 0.0019 (vanilla) and 0.0064 ± 0.0032 (robust).

Nearest Neighbours. Local robustness induces a strong bias towards nearest neighbours, by definition. This assumption is critical as robust training underestimates the effect of making a model robust, treating all the dimen-

	Train	FCs		CNNs	
Shallow Negation <i>(Our benchmark, Barnes et al., 2019)</i>	<i>Vanilla</i>	0.4034 ± 0.0214	0.6303 ± 0.0231	0.3753 ± 0.0091	0.4553 ± 0.0719
	<i>IBP</i> ($\epsilon = 0.001$)	0.3852 ± 0.0071	0.6461 ± 0.0039	0.4954 ± 0.0273*	0.5079 ± 0.0822
	<i>IBP</i> ($\epsilon = 0.01$)	0.4249 ± 0.0260	0.6145 ± 0.0263	0.4715 ± 0.0134*	0.4320 ± 0.0501
Mixed Sentiment <i>(Our benchmark, Barnes et al., 2019)</i>	<i>Vanilla</i>	0.4707 ± 0.0360*	0.6976 ± 0.0126	0.4764 ± 0.0327	0.5506 ± 0.1476
	<i>IBP</i> ($\epsilon = 0.001$)	0.2918 ± 0.0121	0.7205 ± 0.0048	0.5402 ± 0.0961	0.4590 ± 0.1205
	<i>IBP</i> ($\epsilon = 0.01$)	0.2824 ± 0.0169	0.7072 ± 0.0133	0.4485 ± 0.0844	0.5506 ± 0.1476
Sarcasm <i>(Our benchmark, Barnes et al., 2019)</i>	<i>Vanilla</i>	0.5136 ± 0.0504	0.7133 ± 0.0156*	0.4799 ± 0.0393*	0.3067 ± 0.2883
	<i>IBP</i> ($\epsilon = 0.001$)	0.4333 ± 0.0092	0.5578 ± 0.0185	0.6352 ± 0.3962	0.5778 ± 0.3564
	<i>IBP</i> ($\epsilon = 0.01$)	0.4406 ± 0.0943	0.5222 ± 0.0995	0.1650 ± 0.1866	0.1593 ± 0.1030

Table 5: Comparison of 20 *IBP-trained robust* models (Gowal et al. 2018) and their *vanilla* counterparts on samples generated through templates on our benchmark (left subcolumn) and samples exhibiting the same linguistic phenomenon from (Barnes, Øvrelid, and Velldal 2019) (right subcolumn): both populations of networks have been trained on the SST-2 dataset. IBP, which we use to train robust models for two different values of ϵ (0.001 and 0.01), cannot ensure robustness to simple semantic rules and in a few cases worsens the performance of the classifier. Symbol *, when present, means that the improved performance (from *vanilla* to *IBP* or vice-versa) is statistically significant. We consider the two architectures (FCs and CNNs) supported by (Gowal et al. 2018).

sions in the embedding as equally important. We hypothesize this causes the deterioration of the performance of robust models in NLP. The induced invariance along any dimension reduces the effectiveness of the embedding representation on cogent syntactic/semantic tasks such as word-sense-disambiguation, polysemy, etc. *Semantic robustness* takes a different approach and is expected not to be robust to nearest neighbours in the embedding space, but rather to perturbations that are generated by the linguistic rules for which they have been robustly trained. For an increasing number of embedding dimensions, semantic robustness does not suffer in principle from the trade-off between the performance on linguistic tasks (Chen et al. 2013) and robustness guarantees (La Malfa et al. 2020).

Conclusions

In this paper we formalise the concept of *semantic robustness*, which generalizes the notion of NLP robustness by explicitly considering the measurement of robustness on cogent linguistic phenomena. We propose a template-based generative test bed to evaluate semantic robustness. We conduct an empirical analysis that demonstrates that, despite being harder to implement, *semantic robustness* provides stronger guarantees for complex linguistic phenomena where models robust in the classical sense fail. In future, we aim to automate, when possible, the generation of semantic test beds, aided by powerful Masked Language Models such as GPT. We further plan to introduce a validation step for the newly generated texts by involving humans to assess the quality of the semantic perturbations (and consequently of the semantic rules). Finally, we aim to extend our analysis to high-order NLP tasks and study the relationship of linguistic phenomena with out-of-distribution and outlier samples.

Acknowledgements

This project has received funding from the European Research Council (ERC) under the European Union’s Horizon 2020 research and innovation programme (FUN2MODEL, grant agreement No. 834115).

References

- Akhtar, N.; and Mian, A. 2018. Threat of adversarial attacks on deep learning in computer vision: A survey. *IEEE Access*, 6: 14410–14430.
- Alzantot, M.; Sharma, Y.; Elgohary, A.; Ho, B.-J.; Srivastava, M.; and Chang, K.-W. 2018. Generating natural language adversarial examples. *arXiv:1804.07998*.
- Barnes, J.; Øvrelid, L.; and Velldal, E. 2019. Sentiment Analysis Is Not Solved! Assessing and Probing Sentiment Classification. In *Proceedings of the 2019 ACL Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, 12–23. Florence, Italy: Association for Computational Linguistics.
- Barnes, J.; Velldal, E.; and Øvrelid, L. 2021. Improving sentiment analysis with multi-task learning of negation. *Natural Language Engineering*, 27(2): 249–269.
- Baroni, M.; Dinu, G.; and Kruszewski, G. 2014. Don’t count, predict! a systematic comparison of context-counting vs. context-predicting semantic vectors. In *Proceedings of the 52nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 238–247.
- Belinkov, Y.; and Bisk, Y. 2018. Synthetic and Natural Noise Both Break Neural Machine Translation. In *International Conference on Learning Representations*.
- Biggio, B.; Corona, I.; Maiorca, D.; Nelson, B.; Šrndić, N.; Laskov, P.; Giacinto, G.; and Roli, F. 2013. Evasion attacks against machine learning at test time. In *Joint European conference on machine learning and knowledge discovery in databases*, 387–402. Springer.
- Bommasani, R.; Hudson, D. A.; Adeli, E.; Altman, R.; Arora, S.; von Arx, S.; Bernstein, M. S.; Bohg, J.; Bosselet, A.; Brunskill, E.; et al. 2021. On the Opportunities and Risks of Foundation Models. *arXiv:2108.07258*.
- Brown, T. B.; Mann, B.; Ryder, N.; Subbiah, M.; Kaplan, J.; Dhariwal, P.; Neelakantan, A.; Shyam, P.; Sastry, G.; Askell, A.; et al. 2020. Language models are few-shot learners. *arXiv:2005.14165*.
- Chen, Y.; Perozzi, B.; Al-Rfou, R.; and Skiena, S. 2013. The expressive power of word embeddings. *arXiv:1301.3226*.

- Cheng, M.; Yi, J.; Zhang, H.; Chen, P.-Y.; and Hsieh, C.-J. 2018. Seq2sick: Evaluating the robustness of sequence-to-sequence models with adversarial examples. *arXiv:1803.01128*.
- Chollet, F.; et al. 2015. keras.
- Devlin, J.; Chang, M.-W.; Lee, K.; and Toutanova, K. 2018. Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv:1810.04805*.
- Dong, X.; Luu, A. T.; Ji, R.; and Liu, H. 2021. Towards robustness against natural language word substitutions. *arXiv:2107.13541*.
- Ebrahimi, J.; Rao, A.; Lowd, D.; and Dou, D. 2017. Hotflip: White-box adversarial examples for text classification. *arXiv:1712.06751*.
- Ettinger, A.; Rao, S.; Daumé III, H.; and Bender, E. M. 2017. Towards Linguistically Generalizable NLP Systems: A Workshop and Shared Task. In *Proceedings of the First Workshop on Building Linguistically Generalizable NLP Systems*, 1–10.
- Feng, S. Y.; Gangal, V.; Wei, J.; Chandar, S.; Vosoughi, S.; Mitamura, T.; and Hovy, E. 2021. A survey of data augmentation approaches for nlp. *arXiv:2105.03075*.
- Gao, J.; Lanchantin, J.; Soffa, M. L.; and Qi, Y. 2018. Black-box generation of adversarial text sequences to evade deep learning classifiers. In *2018 IEEE Security and Privacy Workshops (SPW)*, 50–56. IEEE.
- Goren, G.; Kurland, O.; Tennenholz, M.; and Raiber, F. 2018. Ranking robustness under adversarial document manipulations. In *The 41st International ACM SIGIR Conference on Research & Development in Information Retrieval*, 395–404.
- Gowal, S.; Dvijotham, K.; Stanforth, R.; Bunel, R.; Qin, C.; Uesato, J.; Arandjelovic, R.; Mann, T.; and Kohli, P. 2018. On the effectiveness of interval bound propagation for training verifiably robust models. *arXiv:1810.12715*.
- Guo, J.; Lu, S.; Cai, H.; Zhang, W.; Yu, Y.; and Wang, J. 2018. Long text generation via adversarial training with leaked information. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 32.
- Hao, Y.; and Veale, T. 2010. An ironic fist in a velvet glove: Creative mis-representation in the construction of ironic similes. *Minds and Machines*, 20(4): 635–650.
- Hendrycks, D.; Liu, X.; Wallace, E.; Dziedzic, A.; Krishnan, R.; and Song, D. 2020. Pretrained transformers improve out-of-distribution robustness. *arXiv preprint arXiv:2004.06100*.
- Hochreiter, S.; and Schmidhuber, J. 1997. Long short-term memory. *Neural computation*, 9(8): 1735–1780.
- Huang, P.-S.; Stanforth, R.; Welbl, J.; Dyer, C.; Yogatama, D.; Gowal, S.; Dvijotham, K.; and Kohli, P. 2019a. Achieving Verified Robustness to Symbol Substitutions via Interval Bound Propagation. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, 4074–4084.
- Huang, P.-S.; Zhang, H.; Jiang, R.; Stanforth, R.; Welbl, J.; Rae, J.; Maini, V.; Yogatama, D.; and Kohli, P. 2019b. Reducing sentiment bias in language models via counterfactual evaluation. *arXiv:1911.03064*.
- Huang, X.; Kwiatkowska, M.; Wang, S.; and Wu, M. 2017. Safety verification of deep neural networks. In *International conference on computer aided verification*, 3–29. Springer.
- Jia, R.; and Liang, P. 2017. Adversarial Examples for Evaluating Reading Comprehension Systems. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, 2021–2031.
- Jia, R.; Raghunathan, A.; Göksel, K.; and Liang, P. 2019. Certified Robustness to Adversarial Word Substitutions. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, 4120–4133.
- Kenyon-Dean, K.; Ahmed, E.; Fujimoto, S.; Georges-Filteau, J.; Glasz, C.; Kaur, B.; Lalande, A.; Bhanderi, S.; Belfer, R.; Kanagasabai, N.; et al. 2018. Sentiment analysis: It’s complicated! In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, 1886–1895.
- Kharitonov, E.; and Chaabouni, R. 2020. What they do when in doubt: a study of inductive biases in seq2seq learners. *arXiv:2006.14953*.
- La Malfa, E.; Wu, M.; Laurenti, L.; Wang, B.; Hartshorn, A.; and Kwiatkowska, M. 2020. Assessing robustness of text classification through maximal safe radius computation. *arXiv:2010.02004*.
- La Malfa, E.; Zbrzezny, A.; Michelmore, R.; Paoletti, N.; and Kwiatkowska, M. 2021. On guaranteed optimal robust explanations for NLP models. *arXiv:2105.03640*.
- Li, J.; Liu, Y.; Chen, T.; Xiao, Z.; Li, Z.; and Wang, J. 2020a. Adversarial attacks and defenses on cyber–physical systems: A survey. *IEEE Internet of Things Journal*, 7(6): 5103–5115.
- Li, L.; Ma, R.; Guo, Q.; Xue, X.; and Qiu, X. 2020b. Bert-attack: Adversarial attack against bert using bert. *arXiv:2004.09984*.
- Liang, B.; Li, H.; Su, M.; Bian, P.; Li, X.; and Shi, W. 2017. Deep text classification can be fooled. *arXiv:1704.08006*.
- Lin, B. Y.; Zhou, W.; Shen, M.; Zhou, P.; Bhagavatula, C.; Choi, Y.; and Ren, X. 2019. CommonGen: A constrained text generation challenge for generative commonsense reasoning. *arXiv:1911.03705*.
- Liu, H.; Dai, Z.; So, D. R.; and Le, Q. V. 2021a. Pay Attention to MLPs. *arXiv preprint arXiv:2105.08050*.
- Liu, J.; Shen, D.; Zhang, Y.; Dolan, B.; Carin, L.; and Chen, W. 2021b. What Makes Good In-Context Examples for GPT-3? *arXiv:2101.06804*.
- Liu, X.; Duh, K.; Liu, L.; and Gao, J. 2020. Very deep transformers for neural machine translation. *arXiv:2008.07772*.
- Liu, Y.; Ott, M.; Goyal, N.; Du, J.; Joshi, M.; Chen, D.; Levy, O.; Lewis, M.; Zettlemoyer, L.; and Stoyanov, V. V.

2019. Roberta: A robustly optimized bert pretraining approach. *arXiv:1907.11692*.
- Madaan, N.; Padhi, I.; Panwar, N.; and Saha, D. 2020. Generate your counterfactuals: Towards controlled counterfactual generation for text. *arXiv:2012.04698*.
- Mikolov, T.; Sutskever, I.; Chen, K.; Corrado, G. S.; and Dean, J. 2013. Distributed representations of words and phrases and their compositionality. In *Advances in neural information processing systems*, 3111–3119.
- Mitchell, T. M. 1980. *The need for biases in learning generalizations*. Department of Computer Science, Laboratory for Computer Science Research
- Morris, J. X. 2020. Second-Order NLP Adversarial Examples. *arXiv:2010.01770*.
- Morris, J. X.; Lifland, E.; Lanchantin, J.; Ji, Y.; and Qi, Y. 2020a. Reevaluating adversarial examples in natural language. *arXiv:2004.14174*.
- Morris, J. X.; Lifland, E.; Yoo, J. Y.; Grigsby, J.; Jin, D.; and Qi, Y. 2020b. Textattack: A framework for adversarial attacks, data augmentation, and adversarial training in nlp. *arXiv:2005.05909*.
- Paszke, A.; Gross, S.; Massa, F.; Lerer, A.; Bradbury, J.; Chanan, G.; Killeen, T.; Lin, Z.; Gimelshein, N.; Antiga, L.; Desmaison, A.; Kopf, A.; Yang, E.; DeVito, Z.; Raison, M.; Tejani, A.; Chilamkurthy, S.; Steiner, B.; Fang, L.; Bai, J.; and Chintala, S. 2019. PyTorch: An Imperative Style, High-Performance Deep Learning Library. In Wallach, H.; Larochelle, H.; Beygelzimer, A.; d'Alché-Buc, F.; Fox, E.; and Garnett, R., eds., *Advances in Neural Information Processing Systems 32*, 8024–8035. Curran Associates, Inc.
- Pennington, J.; Socher, R.; and Manning, C. 2020. GloVe: Global Vectors for Word Representation. In <https://nlp.stanford.edu/projects/glove/>.
- Pröllochs, N.; Feuerriegel, S.; and Neumann, D. 2015. Enhancing sentiment analysis of financial news by detecting negation scopes. In *2015 48th Hawaii International Conference on System Sciences*, 959–968. IEEE.
- Ribeiro, M. T.; Singh, S.; and Guestrin, C. 2018. Semantically equivalent adversarial rules for debugging nlp models. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 856–865.
- Ribeiro, M. T.; Wu, T.; Guestrin, C.; and Singh, S. 2020. Beyond accuracy: Behavioral testing of NLP models with CheckList. *arXiv:2005.04118*.
- Rogers, A.; Kovaleva, O.; and Rumshisky, A. 2020. A primer in bertology: What we know about how bert works. *Transactions of the Association for Computational Linguistics*, 8: 842–866.
- Sinha, K.; Jia, R.; Hupkes, D.; Pineau, J.; Williams, A.; and Kiela, D. 2021. Masked language modeling and the distributional hypothesis: Order word matters pre-training for little. *arXiv:2104.06644*.
- Socher, R.; Perelygin, A.; Wu, J.; Chuang, J.; Manning, C. D.; Ng, A. Y.; and Potts, C. 2013. Recursive deep models for semantic compositionality over a sentiment treebank. In *Proceedings of the 2013 conference on empirical methods in natural language processing*, 1631–1642.
- Song, L.; Yu, X.; Peng, H.-T.; and Narasimhan, K. 2020. Universal adversarial attacks with natural triggers for text classification. *arXiv:2005.00174*.
- Sun, L.; Hashimoto, K.; Yin, W.; Asai, A.; Li, J.; Yu, P.; and Xiong, C. 2020. Adv-BERT: BERT is not robust on misspellings! Generating nature adversarial samples on BERT. *arXiv:2003.04985*.
- Szegedy, C.; Zaremba, W.; Sutskever, I.; Bruna, J.; Erhan, D.; Goodfellow, I.; and Fergus, R. 2014. Intriguing properties of neural networks. In *International Conference on Learning Representations*.
- Tsipras, D.; Santurkar, S.; Engstrom, L.; Turner, A.; and Madry, A. 2018. Robustness may be at odds with accuracy. *arXiv:1805.12152*.
- Vaswani, A.; Shazeer, N.; Parmar, N.; Uszkoreit, J.; Jones, L.; Gomez, A. N.; Kaiser, Ł.; and Polosukhin, I. 2017. Attention is all you need. In *Advances in neural information processing systems*, 5998–6008.
- Wang, R.; Shivanna, R.; Cheng, D.; Jain, S.; Lin, D.; Hong, L.; and Chi, E. 2021. DCN V2: Improved Deep & Cross Network and Practical Lessons for Web-scale Learning to Rank Systems. In *Proceedings of the Web Conference 2021*, 1785–1797.
- Wiegand, M.; Balahur, A.; Roth, B.; Klakow, D.; and Montoyo, A. 2010. A survey on the role of negation in sentiment analysis. In *Proceedings of the workshop on negation and speculation in natural language processing*, 60–68.
- Wu, T.; Ribeiro, M. T.; Heer, J.; and Weld, D. S. 2021. Polyjuice: Automated, general-purpose counterfactual generation. *arXiv:2101.00288*.
- Xu, Y.; Zhong, X.; Yipes, A. J. J.; and Lau, J. H. 2020. Elephant in the room: An evaluation framework for assessing adversarial examples in nlp. *arXiv:2001.07820*.
- Zhang, W. E.; Sheng, Q. Z.; Alhazmi, A.; and Li, C. 2020. Adversarial attacks on deep-learning models in natural language processing: A survey. *ACM Transactions on Intelligent Systems and Technology (TIST)*, 11(3): 1–41.
- Zhang, X.; Zhao, J.; and LeCun, Y. 2015. Character-level convolutional networks for text classification. In *Advances in neural information processing systems*, 649–657.
- Zipf, G. K. 2013. *The psycho-biology of language: An introduction to dynamic philology*. Routledge.